



Training and
Certification

**AWS Technical Essentials
Student Guide
Version 4.2**

100-ESS-42-EN-SG

© 2017 Amazon Web Services, Inc. or its affiliates. All rights reserved.

This work may not be reproduced or redistributed, in whole or in part, without prior written permission from Amazon Web Services, Inc. Commercial copying, lending, or selling is prohibited.

Corrections or feedback on the course, please email us at:

aws-course-feedback@amazon.com.

For all other questions, contact us at:

<https://aws.amazon.com/contact-us/aws-training/>.

All trademarks are the property of their owners.

Contents

Module 1: Introduction and History of AWS	4
Module 2: Foundational Services	29
Module 3: Security, Identity, and Access Management	105
Module 4: Databases	138
Module 5: AWS Elasticity and Management Tools	172
Module 6: Course Wrap-Up	209
Module 7: Course Appendix	222
Module 8: ILT-Specific Content	291

AWS Technical Essentials

Getting Started on AWS

Version 4.2



1

This instructor-led training introduces AWS products and services with exercises and hands-on activities. This training is designed for anyone new to Amazon Web Services so you can gain proficiency in AWS services and make informed decisions about IT solutions based on your business requirements. This training helps you gain the knowledge required to start working on your projects using AWS.

Introductions and Logistics

- Class introductions
- Logistics, bathrooms, breaks
- Participation
- Parking lot
- Cell phones
- Course Errata



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

2

Student errata for the course can be found here: <http://bit.ly/2mZUPPY>

Course Objectives

This course teaches you how to:

- Recognize terminology and concepts as they relate to the **AWS platform** and navigate the **AWS Management Console**.
- Understand the **foundational services**, including Amazon Elastic Compute Cloud (EC2), Amazon Virtual Private Cloud (VPC), Amazon Simple Storage Service (S3), and Amazon Elastic Block Store (EBS).
- Understand the **security** measures AWS provides and key concepts of AWS Identity and Access Management (IAM).
- Understand AWS **database** services, including Amazon DynamoDB and Amazon Relational Database Service (RDS).
- Understand AWS **management tools**, including Auto Scaling, Amazon CloudWatch, Elastic Load Balancing (ELB), and AWS Trusted Advisor.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



3

The course objectives are shown on the slide.

Module Layout

- Module 1: **Introduction and History of AWS**
- Module 2: **Foundational Services** – Amazon EC2, Amazon VPC, Amazon S3, Amazon EBS
 - **Lab 1: Build your VPC and launch a web server**
- Module 3: **Security, Identity, and Access Management** - IAM
- Module 4: **Databases** – Amazon DynamoDB and Amazon RDS
 - **Lab 2: Build your database server and interact with it**
- Module 5: **AWS Elasticity and Management Tools** – Auto Scaling, Elastic Load Balancing, Amazon CloudWatch, and AWS Trusted Advisor
 - **Lab 3: Scale and load balance your application and monitor activity**
- Module 6: Course Wrap-Up
- Module 7: Course Appendix

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



4

Module 1

Introduction and History of AWS




5

Module 1 describes the history and fundamental elements of Amazon Web Services (AWS), and shows you how to navigate the AWS Management Console. The AWS Management Console is the graphical user interface (GUI) to access AWS. AWS is also accessible through the command line interface (CLI) and software development kits (SDKs). This module discusses the AWS global infrastructure and basic principles of deploying on AWS.


By the end of this module you will be able to:

- Describe the history of AWS.
- Recognize the AWS global infrastructure.
- Navigate the AWS Management Console.



Jeff Bezos incorporated the company in 1994, and Amazon.com was launched in 1995 as an online bookstore. Amazon.com, Inc. is an American multinational electronic commerce company with its headquarters in Seattle, Washington. It is the world's largest online retailer. Amazon has continued to grow and officially launched Amazon Web Services (AWS) in 2006. More came after, including Amazon Publishing, the Kindle, Amazon Game Studios, and Amazon Art.

After over a decade of building and running the highly scalable web application Amazon.com, the company realized that it had developed a core competency in operating massive-scale technology infrastructure and data centers and embarked on a much broader mission of serving a new customer segment—developers and businesses—with a platform of web services they can use to build sophisticated, scalable applications. Today, AWS is the fastest-growing multi billion-dollar enterprise IT vendor in the world.




Amazon Web Services is 10+ years in the making. Amazon Web Services, also abbreviated to AWS, is a collection of remote computing services called *web services*. These web services make up a cloud computing platform offered via the Internet. We deliver web-based cloud services for storage, computing, networking, databases, and more.

The AWS mission is to enable businesses and developers to use web services to build scalable, sophisticated applications. Web services is another name for what people now call “the cloud.”

For more information, see: <http://aws.amazon.com>

AWS Pace of Innovation

AWS has been continually expanding its services to support virtually any cloud workload, and it now has more than 90 services that range from compute, storage, networking, database, analytics, application services, deployment, management, developer, mobile, Internet of Things (IoT), Artificial Intelligence (AI), security, hybrid and enterprise applications. AWS has launched a total of 1,017 new features and/or services year to date* - for a total of 2,913 new features and/or services since inception in 2006.



* As of 1 January 2017




8

AWS is continually expanding its services to support virtually any cloud workload.

Innovation is in our DNA, and our structure and approach to product development and delivery is fundamentally different from that of other IT vendors. We have decentralized, autonomous development teams that are working directly with customers. They are empowered to develop and launch based on what they learn from interactions with customers. We iterate products continuously, and the newest is instantly available to customers. No need to upgrade, deploy, or migrate. When a feature or enhancement is ready, we “push” it out.

This approach also enables us to very rapidly introduce and iterate on new services.

For more information, see: <https://aws.amazon.com/blogs/aws/>



As of January 1, 2017, AWS has launched 2,913 new services and updates to existing services.

AWS Customers

Enterprise Customers

Amazon Web Services delivers a mature set of services specifically designed for the unique security, compliance, privacy, and governance requirements of large organizations.



Public Sector

Paving the way for innovation and supporting world-changing projects in government, education and nonprofit organizations.



Startups

From the spark of an idea, to your first customer, to IPO and beyond, let Amazon Web Services help you build and grow your startup.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

10

Enterprise Customers: Enterprise cloud computing with AWS can help IT increase innovation, agility, and resiliency; all while reducing cost. With AWS, you can build enterprise cloud solutions quickly and without a big up-front investment. The free tier allows you to prototype virtually any application for free.

Startup Customers: Our innovations free you to scale quickly, go to market faster, and control costs. AWS Activate is a free program with resources for startups to get the most out of AWS from day one.

Public Sector Customers: AWS offers scalable, cost-effective cloud services that public sector customers can use to meet mandates, reduce costs, drive efficiencies, and accelerate innovation.

For more information, see:

- <http://aws.amazon.com/contract-center/>
- Enterprise Cloud Computing: <http://aws.amazon.com/enterprise/>
- Free Tier: <https://aws.amazon.com/free/>
- Startups on Amazon Web Services: <http://aws.amazon.com/start-ups/>

- AWS in the Public Sector: <https://aws.amazon.com/government-education/>


Advantages and Benefits of AWS Cloud Computing



Trade capital expense
for variable expense.



Benefit from massive
economies of scale.



Stop guessing
capacity.



Increase speed and
agility.



Stop spending money on
running and maintaining
data centers.



Go global in minutes.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

11

Cloud computing provides a simple way to access servers, storage, databases, and a broad set of application services over the Internet. AWS owns and maintains the network-connected hardware required for these application services, while you provision and use what you need via a web application.

Trade capital expense for variable expense: Instead of investing heavily in data centers and servers before you know how you're going to use them, you can pay only when you consume computing resources and pay only the amount you consume.

Benefit from massive economies of scale: By using cloud computing, you can achieve a lower variable cost than you can get on your own. Because usage from hundreds of thousands of customers is aggregated in the cloud, providers such as Amazon Web Services can achieve higher economies of scale, which translates into lower pay-as-you-go prices.

Stop guessing capacity: Eliminate guessing about your infrastructure capacity needs. When you make a capacity decision before deploying an application, you often end up either sitting on expensive idle resources or dealing with limited capacity. With cloud computing, these problems go away. You can access as much or as little as you need and scale up and down as required with only a few minutes' notice.

Increase speed and agility: In a cloud computing environment, new IT resources are only ever a click away, which means you reduce the time it takes to make those resources available to your developers from weeks to just minutes. This results in a dramatic increase in agility for the organization, because the cost and time it takes to experiment and develop is significantly lower.

Stop spending money on running and maintaining data centers: Focus on projects that differentiate your business, not the infrastructure. Cloud computing lets you focus on your own customers, instead of on the heavy lifting of racking, stacking, and powering servers.

Go global in minutes: Easily deploy your application in multiple regions around the world with just a few clicks. This means you can provide a lower latency and better experience for your customers simply and at minimal cost.

AWS Positioned as a Leader in the Gartner Magic Quadrant for Cloud Infrastructure as a Service, Worldwide*

AWS is positioned highest in execution and furthest in vision within the Leaders Quadrant



Source: Gartner (August 2016)

*Gartner, Magic Quadrant for Cloud Infrastructure as a Service, Worldwide, Leong, Lydia, Petri, Gregor, Gill, Bob, Dorosh, Mike, August 32016


This graphic was published by Gartner, Inc. as part of a larger research document and should be evaluated in the context of the entire document. The Gartner document is available upon request from AWS.

Gartner does not endorse any vendor, product or service detailed in its research publications, and does not advise technology users to select only those vendors with the highest ratings or other designation. Gartner research publications consist of the opinions of Gartner's research organization and should not be construed as statements of fact. Gartner disclaims all warranties, expressed or implied, with respect to this research, including any warranties of merchantability or fitness for a particular purpose.




Training and
Certification

12



Many of our services and terminology have analogs in the traditional IT space. This side-by-side comparison shows how AWS products and services relate to a traditional infrastructure.



AWS cloud computing provides a simple way to access servers, storage, databases, and a broad set of application services over the Internet. AWS owns and maintains the network-connected hardware required for these application services, while you provision and use what you need.

AWS Foundation Services

Compute

- AmazonEC2
- AmazonEC2 Container Registry
- AmazonEC2 Container Service
- Amazon Lightsail
- Amazon VPC
- AWS Batch
- AWS Elastic Beanstalk
- AWS Lambda
- Elastic Load Balancing

Network

- Amazon CloudFront
- Amazon Route 53
- Amazon VPC
- AWS Direct Connect
- Elastic Load Balancing

Storage

- Amazon EFS
- Amazon Glacier
- Amazon S3
- AWS Snowball
- AWS Storage Gateway

Security & Identity

- Amazon Inspector
- AWS Artifact
- AWS Certificate Manager
- AWS CloudHSM
- AWS Directory Service
- IAM
- AWS KMS
- AWS Organizations
- AWS Shield
- AWS WAF

Applications

- Amazon WorkDocs
- Amazon WorkMail
- Amazon AppStream
- Amazon WorkSpaces

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

15

AWS foundation services are categorized as shown in the slide.

AWS Platform Services

Databases	Analytics	Application Services	Management Tools	Developer Tools	Mobile Services	Internet of Things
Amazon DynamoDB	Amazon Athena	Amazon API Gateway	Amazon CloudWatch	AWS CodeBuild	Amazon API Gateway	AWS IoT
Amazon ElastiCache	Amazon CloudSearch	Amazon AppStream 2.0	AWS CloudFormation	AWS CodeCommit	Amazon Cognito	AWS Greengrass
Amazon RDS	Amazon EMR	Amazon Elastic Transcoder	AWS CloudTrail	AWS CodeDeploy	Amazon Mobile Analytics	
Amazon Redshift	Amazon ES	Amazon SWF	AWS Config	AWS CodePipeline	Amazon Pinpoint	
	Amazon Kinesis	AWS Step Functions	AWS Managed Services	AWS X-Ray	AWS Device Farm	
	Amazon QuickSight		AWS OpsWorks		AWS Mobile Hub	
	Amazon Redshift		AWS Service Catalog			
			AWS Trusted Advisor			

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

16

AWS foundation services are categorized as shown in the slide.

AWS Global Infrastructure

Regions

- Geographic locations
- Consist of **at least two** Availability Zones

Availability Zones


- Clusters of data centers
- **Isolated from failures** in other Availability Zones

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



17

AWS Regions are geographic locations that contain multiple Availability Zones. Availability Zones consist of data centers clustered in a region. Each Availability Zone is engineered to be isolated from failures in other Availability Zones.



AWS is steadily expanding its global infrastructure to help customers achieve lower latency and higher throughput and to ensure that your data resides only in the region you specify. As you and all customers grow their businesses, AWS will continue to provide infrastructure that meets your global requirements.

As of March 2017, the AWS Cloud operates 42 Availability Zones within 16 geographic Regions around the world, with five more Availability Zones and two more Regions coming online throughout the next year.

The AWS GovCloud (US) Region is an isolated region designed to allow US government agencies and customers to move sensitive workloads into the cloud by addressing their specific regulatory and compliance requirements. AWS products and services are available by region so you may not see all regions available for a given service.

You can run applications and workloads from a region to reduce latency to end users while avoiding the up-front expenses, long-term commitments, and scaling challenges associated with maintaining and operating a global infrastructure.

For more information, see: <http://aws.amazon.com/about-aws/global-infrastructure/>

AWS Global Infrastructure


At least 2 Availability Zones per region.

Examples:

- US East (N. Virginia)
 - us-east-1a
 - us-east-1b
 - us-east-1c
 - us-east-1d
 - us-east-1e



- Asia Pacific (Tokyo)
 - ap-northeast-1a
 - ap-northeast-1b
 - ap-northeast-1c



Note: Conceptual drawing only. The number of Availability Zones (AZ) may vary.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification


19

Each region is a separate geographic area that has multiple locations isolated from each other, known as Availability Zones. Each Availability Zone is isolated, but the Availability Zones in a region are connected through low-latency links. Where natural disasters or fault lines are a consideration, AWS isolates its Availability Zones so that they are not easily affected at the same time. For example, where earthquakes are a problem, AWS would not build two Availability Zones on the same fault line. When you launch an instance, you can select an Availability Zone or let AWS choose one for you. If you distribute your instances across multiple Availability Zones, you can design your application so that if an instance fails, an instance in another Availability Zone can handle requests.

For more information, see:

- <http://aws.amazon.com/about-aws/global-infrastructure/regional-product-services/>
- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/using-regions-availability-zones.html>
- <https://www.amazonaws.cn/en/>

High Availability Using Multi-AZ Deployments



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

20

AWS highly recommends that you provision your compute resources across multiple Availability Zones. If you have multiple instances, you can run them across more than one Availability Zone and get added redundancy. If a single Availability Zone has a problem, all assets in your second Availability Zone are unaffected.

AWS Global Infrastructure – Edge Locations

- 70* edge locations
- Local points of presence that support AWS services like:

 **Amazon Route 53**

 **Amazon CloudFront**

 **AWS WAF**

 **AWS Shield**

*as of March 2017

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



21

Edge locations help lower latency and improve performance for end users.

AWS Management Console Demonstration



Training and
Certification

22

Look at the AWS Management Console so you are familiar with the navigation:

- Selecting a region.
- Dragging and dropping service icons at the top for quick use.
- Verifying what is included in the left navigation pane.

For more information, see: <https://console.aws.amazon.com/console/home>

Knowledge Check

Q: What is the AWS term for physically distinct groups of **data centers** within a region?

Availability Zone

True or False: There are more **Regions** than **Edge locations**.

False

True or False: AWS owns and maintains the infrastructure required for application services. You provision and use them as needed.

True

Q: How do **Availability Zones** in the same region differ?

Each Availability Zone is isolated, but the Availability Zones in a region are connected through low-latency links.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



23

In this module, you learned about the history of Amazon and AWS, the global infrastructure of AWS, and how to navigate the AWS Management Console.

Test out some of your new skills!

Module 2

AWS Foundational Services

Module 2 Layout

- Amazon Elastic Compute Cloud (EC2)
- Amazon Virtual Private Cloud (VPC)
 - **Lab 1: Build your VPC and launch a web server**
- Amazon Storage Services
 - Amazon Simple Storage Service (S3)
 - Amazon Elastic Block Store (EBS)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



2

Amazon Elastic Compute Cloud (EC2)



3

Understand Amazon Elastic Compute Cloud (EC2) concepts, including:

- Instances vs. servers
- Types and families
- Ephemeral vs. persistent storage (root instance volumes)
- Amazon Machine Images (AMIs)
- Bootstrapping/user data

Amazon Elastic Compute Cloud (EC2)



Amazon
EC2

- **Resizable** compute capacity
- Complete control of your computing resources
- **Reduced time required** to obtain and boot new server instances

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



4

Amazon EC2 instances are virtualized servers in Amazon's data centers.

Amazon EC2 is designed to make web-scale computing easier for developers. With Amazon EC2's simple web service interface, you can obtain and configure capacity with minimal friction. It provides you with complete control of your computing resources and allows you to run on Amazon's proven computing environment.

Amazon EC2 reduces the time required to obtain and boot new server instances, which helps you to quickly scale capacity as your computing requirements change. Amazon EC2 changes the economics of computing by allowing you to pay only for capacity that you actually use. Amazon EC2 provides the tools for you to build failure-resilient applications and isolate the applications from common failure scenarios.

For more information, see: <https://aws.amazon.com/ec2/>

Amazon EC2 Facts

- **Scale capacity** as your computing requirements change
- Pay only for capacity that you actually use
- Choose **Linux or Windows**
- Deploy across **AWS Regions** and **Availability Zones** for reliability
- Use **tags** to help manage your Amazon EC2 resources



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

5

Amazon EC2 presents a true virtual computing environment, allowing you to use web service interfaces to launch instances with a variety of operating systems, load them with your custom application environment, manage your network's access permissions, and run your image using as many systems as you need.

You can programmatically scale your computing capacity as your requirements change. You pay only for capacity that you actually use and can choose Linux or Windows. You can leverage the AWS global infrastructure to deploy across regions and Availability Zones (AZs) for reliability.

For more information, see:

http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/Using_Tags.html

Launching an Amazon EC2 Instance via the Management Console



- 1. Determine the AWS Region** in which you want to launch the Amazon EC2 instance.
- 2. Launch** an Amazon EC2 instance from a pre-configured Amazon Machine Image (AMI).
- 3. Choose an instance type** based on CPU, memory, storage, and network requirements.
- 4. Configure** network, IP address, security groups, storage volume, tags, and key pair.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



6

Before you create your first Amazon EC2 instance, decide which region you want to have that instance in. The AMI comes pre-installed with many AWS API tools and with CloudInit. AWS API tools enable scripting of important provisioning tasks from within an Amazon EC2 instance. AMIs are like building blocks of EC2 instances. They are templates of a computer's volumes. AMIs can have public or private access. You can also create gold master images of your Amazon EC2 infrastructure, which allow you to decrease your boot times.

Amazon Machine Image (AMI) Details



An AMI includes the following:

- A template for the **root volume** for the instance (for example, an operating system, an application server, and applications).
- **Launch permissions** that control which AWS accounts can use the AMI to launch instances.
- A block device mapping that specifies the **volumes to attach** to the instance when it is launched.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

7

An AMI is a template that contains a software configuration such as an operating system, application server, and applications. You use an AMI to launch an instance, which is the copy of the AMI running as a virtual server on a host computer in Amazon's data center. You can launch as many instances as you want from an AMI. You can also launch instances from as many AMIs as you need.

You can create your own AMI by customizing the instance that you launch from a public AMI and then saving the configuration as a custom AMI for your own use. You can also buy, share, and sell AMIs.


For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/AMIs.html>

Instances and AMIs

Select an AMI based on:

- Region
- Operating system
- Architecture (32-bit or 64-bit)
- Launch permissions
- Storage for the root device




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

8

You select an AMI based on region, operating system, architecture, launch permissions, and storage for the root device. Launch permissions determine availability of an AMI and are either public (the owner grants launch permissions to all AWS accounts), explicit (the owner grants launch permissions to specific AWS accounts), or implicit (the owner has implicit launch permissions for an AMI).




You can launch multiple instances of different types from a single AMI when launching an EC2 instance. An instance type essentially determines the hardware of the host computer used for your instance. Each instance type offers different compute and memory capabilities. Select an instance type based on the amount of memory and computing power that you need for the application or software that you plan to run on the instance. Your instance keeps running until you stop or terminate it or until it fails.

Instances are deployed in the Amazon EC2 public cloud or the Amazon Virtual Private Cloud in an Availability Zone within a region. You can configure security and network access on your Amazon EC2 instance.

Customers can deploy to multiple Availability Zones within a region. You choose which instance types you want, and then start, terminate, and monitor as many instances of your AMI as needed, using the web service APIs or the variety of management tools provided.

Amazon EC2 instances can leverage Amazon Elastic Block Store (EBS) volumes in each Availability Zone. Determine whether you want to run in multiple locations, use static IP endpoints, or attach persistent block storage to your instances. Amazon EBS volumes can be saved using “snapshots.” Additionally, Amazon S3 buckets can be used to store data objects needed by Amazon EC2 instances. Pay only for the resources that you actually consume, such as instance-hours or data transfer.



The slide shows the lifecycle of an instance launched from an AMI. Note that you can only stop and start instances that are EBS-backed.

An EC2 instance can be in one of the following states:

- Pending: When you launch an instance, it enters the pending state and the instance moves to a new host computer. The instance type specified at launch determines the hardware of the host computer for your instance.
- Running: AWS uses the AMI specified at launch to boot the instance. When the instance is ready for you, it enters the running state. You can connect to your running instance and use it as you would a computer sitting in front of you. When your instance is in the running state, you're billed for each hour or partial hour that you keep the instance running. You are billed for all running instances, even if they are idle and not being connected to.
- Rebooting: You can reboot your instance through the Amazon EC2 console, Amazon EC2 CLI, and the Amazon EC2 API. It is recommended that you reboot your EC2 instance instead of running the operating system reboot from the instance. When an instance is rebooted, it remains on the same host computer and maintains its public DNS name, private IP address, and any data on its instance store volumes. Rebooting an instance doesn't start a new instance billing hour.


- Shutting down: When you've decided that you no longer need an instance, you can terminate it. The instance will enter the shutting-down state. You will stop incurring charges as soon as the instance enters shutting-down or terminated states.
- Terminated: A terminated instance remains visible in the console for a while before it is deleted. You cannot connect to or recover a terminated instance.
- Stopping: Amazon EBS-backed instances can be stopped. When you stop an instance, it enters the stopping state.
- Stopped: Amazon EBS-backed instances in the stopped state are no longer eligible for hourly usage or data transfer fees. AWS does charge for the storage of EBS volumes on stopped instances. You can modify certain attributes of stopped instances, including the instance type. Starting a stopped instance puts it back into the pending state, which moves the instance to a new host machine. When you stop and start an instance, you lose any data on the instance store volumes on the previous host computer.

AWS Marketplace – IT Software Optimized for the Cloud

- Online store to discover, purchase, and deploy IT software on top of the AWS infrastructure.
- Catalog of **2700+** IT software solutions including Paid, BYOL, Open Source, SaaS, and free-to-try options.
- Pre-configured to operate on AWS.
- Software checked by AWS for security and operability.
- Deploys to AWS environment in minutes.
- Flexible, usage-based billing models.
- Software charges billed to AWS account.

Includes [AWS Test Drive](#).

<https://aws.amazon.com/marketplace>



AWS and Oracle have worked together to offer customers convenient options for deploying enterprise applications in the cloud. Customers can build enterprise-grade solutions hosted by Amazon Web Services using database and middleware software by Oracle, and can also launch entire enterprise software stacks from Oracle on Amazon EC2.

New and existing SAP customers can confidently deploy their SAP solutions on SAP-certified Amazon EC2 instances in production environments because SAP and AWS have tested the performance of the underlying AWS resources, verified their performance, and certified them against the same standards that apply to servers and virtual platforms.

AWS also provides infrastructure services that allow customers to easily run Microsoft Windows Server applications in the cloud, without the cost and complexity of having to purchase or manage servers or data centers. Available AMIs allow customers to start running fully supported Windows Server virtual machine instances in minutes.

Customers may also rely on the global infrastructure of AWS to power everything from custom .NET applications to enterprise deployments of Microsoft Exchange Server, SQL Server, or SharePoint Server.

Software launched from AWS customers automatically deploys onto Amazon Elastic Compute Cloud (EC2), which is the AWS compute service. AWS customers use 143 million hours a month of Amazon EC2 for AWS Marketplace software products.

The benefits of AWS Marketplace include:

- Easy product discovery
- Streamlined buying experience
- Simplified billing
- Expedited deployment cycles
- Optimized software capacity
- Matched spend to actual usage
- Trust vetted and scanned products

AWS Test Drive provides a private IT sandbox environment that contains preconfigured server based solutions. In under an hour, and using a step-by-step lab manual and video, you can launch, sign in, and learn about these popular third-party IT solutions powered by AWS and AWS CloudFormation.

For more information, see: <https://aws.amazon.com/marketplace>

Choosing the Right Amazon EC2 Instance



AWS uses Intel® Xeon® processors to provide customers with high performance and value. EC2 instance types are optimized for different use cases, workload requirements and come in multiple sizes.

Consider the following when choosing your instances:

- Core count
- Memory size
- Storage size and type
- Network performance
- CPU technologies

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



12

Choosing the right EC2 instance type matters.

AWS has a wide variety of EC2 compute instances. Most are based on Intel Xeon processors.

Intel® Xeon processors are compatible with the broadest range of ecosystem software, which helps developers to develop their application rapidly. IT also values Xeon processors for higher performance and reliability.

Each instance type or family is optimized for different workloads or use cases. Within each type or family we have multiple sizes: Large, XLarge, 2XLarge, etc.

When you choose your instance type, you should consider the several different attributes of each family, such as number of cores, amount of memory, amount and type of storage, network performance, and Intel processor technologies.

Larger instances are better for workloads that scale.

Get the Intel® Advantage



Intel's Haswell microarchitecture on new X1, C4, D2, and M4 instances, with **custom Intel® Xeon® v3** processors, provides new features:

Haswell microarchitecture can boost existing applications performance by **30% or more** for better workload performance and faster response times.

Hardware Assisted technologies, such as **Intel® AVX2.0** instructions, can double the floating-point performance for compute-intensive workloads and provide additional instructions for compression and encryption

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon Web Services | Training and Certification

13

In 2016, AWS launched our C4, D2, and M4 instances based on Intel's latest 22nm Haswell microarchitecture. They all use custom Intel® Xeon® v3 processors designed and built especially for AWS.

X1 is based on Intel E7 enterprise class CPU.

Haswell can boost application performance, which helps workloads to perform better or deliver faster response time. Haswell instances also include newer hardware-assisted technologies for workload acceleration, such as AVX 2.0, which doubles the floating point performance compared to AVX 1.0 for high performance compute (HPC) and financial modeling kinds of workload.

X1 Instance - Tons of Memory

The X1 instance:

- Features up to 2TB of memory and 100 vCPU.
- Uses Intel E7 v3 Haswell processors.
- Is designed for demanding enterprise workloads, including production installations of SAP HANA, Microsoft SQL Server, Apache Spark, and Presto.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon Web Services | Training and Certification

14

For more information, see: <https://aws.amazon.com/blogs/aws/ec2-instance-update-x1-sap-hana-t2-nano-websites/>

Intel® Processor Technologies

Intel® AVX: Provides dramatically better performance for highly parallel HPC workloads such as *life science engineering, data mining, financial analysis*, or other technical computing applications. AVX also enhances *image, video, and audio processing*.

Intel® AES-NI: Enhance your security with these new encryption instructions that reduce the performance penalty associated with encrypting/decrypting data.

Intel® Turbo Boost Technology: Provides more computing power when you need it with performance that adapts to spikes in your workload.

Intel Transactional Synchronization (TSX) Extensions: Enable execution of transactions that are independent to accelerate throughput.

P state & C state control: Gives you the ability to individually tune each cores performance & sleep states to improve application performance.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



15

In addition to the new X1, C4, D2, and M4 instance families based on Haswell, with its host of new microarchitecture features, our other EC2 instances are based on Intel Xeon processors. These Xeon processors have important technology features that you should be aware of.

Intel® AVX is perfect for highly parallel HPC workloads, such as life sciences or financial analysis.

Intel® AES-NI accelerates encryption/decryption of data and therefore reduces the penalty you would pay by using encryption.

Intel® Turbo Boost Technology automatically gives you more computing power when your workloads are not fully using all CPU cores. Think of it as automatic overclocking when you have thermal headroom.

Intel Transactional Synchronization (TSX) Extensions enable execution of transactions that are independent in order to accelerate throughput. Intel® TSX delivers outstanding performance gains for multi-threaded applications that are developed using a coarse-grained lock (a single lock for the entire hash table) that behaves like a fine-grained lock (multiple locks for smaller table sections). Coarse-grained locks are easier to use, easier to understand, and easier to debug. In summary, Intel TSX provides a set of instruction set extensions that allow programmers to specify regions of code for transactional synchronization. Programmers can use these extensions to achieve the performance of fine-grained locking while actually programming using coarse-grained locks.

P state and C state control on the 8xlarge (when a user has the whole 2S machine) provides the ability to individually tune each core's performance and sleep states to improve application performance.

AWS EC2 Instances with Intel® Xeon®

AWS Instance Type	High Memory X1	Compute-Optimized C4	Storage-Optimized D2	General Purpose M4	Memory-Optimized R3	IO-Optimized I3	Graphics-Optimized G2	Burstable Performance T2
Intel Xeon Processor	E7-8880 v3	E5-2666 v3	E5-2676 v3	E5-2686 v4 E5-2676 v3	E5-2670 v2	E5-2686 v4	E5-2670	Intel Xeon Family
Intel AVX	AVX 2.0	AVX 2.0	AVX 2.0	AVX 2.0	Yes	Yes	Yes	Yes
Intel AES-NI	Yes	Yes	Yes	Yes	Yes	Yes	No	No
Intel Turbo Boost	Yes	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Intel TSX	Yes	No	No	No	No	No	No	No
Per core P- and C-state control	No	Yes (8xlarge only)	No	No	No	No	No	No
SSD Storage	EBS Optimized by default	EBS Optimized by default	No	EBS: Optimized by default	Yes	Yes	Yes	EBS only

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



16

This chart provides you with a quick reference to the individual Intel® technologies that were explained on the previous slide and to which EC2 instance contains each Intel technology.

For each instance, a few key technologies are identified to end customers because they have significant impact on the performance of some workloads. More details are discussed later in the presentation, but for now, be aware that different instances have different instructions and capabilities that will affect performance beyond just frequency.

Current Generation Instances



Instance Family	Some Use Cases
General purpose (t2, m4, m3)	<ul style="list-style-type: none">Low-traffic websites and web applicationsSmall databases and mid-size databases
Compute-optimized (c4, c3)	<ul style="list-style-type: none">High performance front-end fleetsVideo-encoding
Memory-optimized (r3)	<ul style="list-style-type: none">High performance databasesDistributed memory caches
Storage-optimized (i2, d2)	<ul style="list-style-type: none">Data warehousingLog or data-processing applications
GPU instances (g2)	<ul style="list-style-type: none">3D application streamingMachine learning

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



17

Each vCPU is a hyperthread of an Intel Xeon core for M4, M3, C4, C3, R3, HS1, G2, I2, and D2.

Amazon EC2 lets you choose from several different instance types to meet your computing needs. Each instance provides a predictable amount of dedicated compute capacity and is charged per instance-hour consumed.

First-generation

(M1) general purpose instances provide a balanced set of resources and a low-cost platform that is well suited for a wide variety of applications.

Second-generation

(M3) general purpose instances provide a balanced set of resources and a higher level of processing performance compared to first-generation general purpose instances. Instances in this family are ideal for applications that require higher absolute CPU and memory performance. Applications that can benefit from the performance of second-generation general purpose instances include encoding applications, high traffic content management systems, and Memcached applications. High-memory instances offer large memory sizes for high-throughput applications, including database and memory-caching applications. High-CPU instances have proportionally more CPU resources than memory (RAM) and are well suited for compute-intensive applications. There are also various high-storage and cluster-computer instance types available.

For more information, see:

- <http://docs.aws.amazon.com/AWSEC2/latest/WindowsGuide/instance-types.html>
- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instance-types.html>
- <http://aws.amazon.com/ec2/instance-types/>

Instance Metadata

- Is **data** about your **instance**.
- Can be used to **configure or manage** a running instance.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

18

Instance metadata is divided into categories.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html#instancedata-data-categories>

Retrieving Instance Metadata

To view all categories of instance metadata from within a running instance, use the following URI:

<http://169.254.169.254/latest/meta-data/>

On a Linux instance, you can use:

- `$ curl http://169.254.169.254/latest/meta-data/`
- `$ GET http://169.254.169.254/latest/meta-data/`

All metadata is returned as text (content type `text/plain`).



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

19

Because your instance metadata is available from your running instance, you do not need to use the Amazon EC2 console or the AWS CLI. This can be helpful when you're writing scripts to run from your instance. Note that you are not billed for HTTP requests that are used to retrieve instance metadata and user data.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html>

Instance User Data



- Can be passed to the instance **at launch**.
- Can be used to perform common **automated configuration tasks**.
- Runs scripts after the instance starts.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification

20

Instance user data can be used to build more generic AMIs that can be modified by configuration files supplied at launch time. It is important to note that, although you can only access instance metadata and user data from within the instance itself, the data is not protected by cryptographic methods. Anyone who can access the instance can view its metadata.

Adding User Data



- You can specify user data when launching an instance.
- User data can be:
 - Linux script – executed by **cloud-init**
 - Windows batch or PowerShell scripts – executed by **EC2Config** service
- User data scripts run once per instance ID by default.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



21

You can specify user data to configure an instance during launch or to run a configuration script. To attach a file, select the As file option and browse for the file to attach. The cloud-init package is an open-source application built by Canonical that is used to bootstrap Linux images in a cloud computing environment, such as Amazon EC2.

User data information:

- User data is treated as opaque data: what you give is what you get back. The instance is responsible for interpreting user data.
- User data is limited to 16 KB. This limit applies to the data in raw form, not base64-encoded form.
- User data must be base64-encoded before being submitted to the API. The Amazon EC2 command line tools perform the base64 encoding for you. The data is decoded before being presented to the instance.
- User data is executed only at launch. If you stop an instance, modify the user data, and start the instance, the new user data is not executed automatically by default.

User Data Example Linux



```
#!/bin/sh
yum -y install httpd
chkconfig httpd on
/etc/init.d/httpd start
```

User data shell scripts must start with the #! characters and the path to the interpreter you want to read the script.

Install Apache web server
Enable the web server
Start the web server

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon | Training and Certification

The slide shows an example of user data on Linux. You can also provide user data to an instance on Linux by using the #cloud-config directive – a format defined by cloud-init.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/user-data.html>

User Data Example Windows



```
<powershell>
```

```
Import-Module ServerManager
```

Import the Server Manager module
for Windows PowerShell.

```
Install-WindowsFeature web-server, web-webserver
```

```
Install-WindowsFeature web-mgmt-tools
```

```
</powershell>
```

Install IIS
Install Web Management Tools

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and
Certification

23

You can send user data to a Windows instance with a PowerShell script (shown in slide) or with a set of Windows batch commands.


Retrieving User Data

To retrieve user data, use the following URL:

<http://169.254.169.254/latest/user-data>

On a Linux instance, you can use:

```
$ curl http://169.254.169.254/latest/user-data/  
$ GET http://169.254.169.254/latest/user-data/
```



The screenshot shows a terminal window titled "ec2-user@ip-172-31-31-72:~". The user is authenticated using the public key "imported-openssh-key". The command entered is "curl http://169.254.169.254/latest/user-data". The terminal output shows the retrieval of the user data file, which contains PHP code to display the PHPinfo() output. The file is saved as "phpinfo.php" in the "/var/www/html/" directory.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Training and
Certification

24

To retrieve instance user data, use the following URI:

<http://169.254.169.254/latest/user-data>

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-metadata.html#instancedata-user-data-retrieval>

Amazon EC2 Purchasing Options



On-Demand Instances Pay by the hour.	Reserved Instances Purchase, at a significant discount instances that are always available 1-year to 3-year terms.	Scheduled Instances Purchase instances that are always available on the specified recurring schedule, for a one-year term.	Spot Instances Bid on unused instances which can run as long as they are available and your bid is above the Spot price.	Dedicated Instances Pay, by the hour, for instances that run on single-tenant hardware.	Dedicated Hosts Pay for a physical host that is fully dedicated to running your instances.
--	---	--	--	---	--

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification 25

On-Demand Instances are free-tier eligible. They have the lowest up-front cost and the most flexibility. You pay for an hour at a time with no up-front commitments or long-term contracts. This is great for applications with short-term, spiky, or unpredictable workloads.

Amazon EC2 Reserved Instance pricing allows you to reserve computing capacity for 1-year to 3-year terms at a significantly discounted hourly rate. Reserved Instances provide a billing discount and capacity reservation that is applied to instances to lower hourly running costs. A Reserved Instance is not a physical instance. The discounted usage price is fixed for as long as you own the Reserved Instance, which allows you to predict compute costs over the term of the reservation. If you are expecting consistent, heavy, use, Reserved Instances can provide substantial savings over owning your own hardware or running only On-Demand Instances.

Scheduled Reserved Instances enable you to purchase capacity reservations that recur on a daily, weekly, or monthly basis, with a specified duration, for a 1-year term. You reserve the capacity in advance, so that you know it is available when you need it. You pay for the time the instances are scheduled, even if you do not use them. Scheduled instances are a good choice for workloads that do not run continuously, but do run on a regular schedule and take a finite time to complete.

Spot Instances enable you to bid on unused EC2 instances, which can lower your Amazon EC2 costs significantly. The hourly price for a Spot instance (of each instance type in each Availability Zone) is set by Amazon EC2 and fluctuates depending on the supply of and demand for Spot instances. Your Spot Instance runs whenever your bid exceeds the current market price.

Spot Instances are a cost-effective choice if you can be flexible about when your applications run and if your applications can be interrupted. Amazon EC2 does not terminate Spot instances with a specified duration (also known as Spot blocks) when the Spot price changes. This makes them ideal for jobs that take a finite time to complete, such as batch processing, encoding and rendering, modeling and analysis, and continuous integration.

An Amazon EC2 Dedicated Host is a physical server with EC2 instance capacity fully dedicated to your use. Dedicated Hosts allow you to use your existing per-socket, per-core, or per-VM software licenses, including Microsoft Windows Server, Microsoft SQL Server, SUSE, and Linux Enterprise Server. Dedicated Hosts and Dedicated Instances can both be used to launch Amazon EC2 instances onto physical servers that are dedicated for your use. There are no performance, security, or physical differences between Dedicated Instances and instances on Dedicated Hosts. However, Dedicated Hosts give you additional visibility and control over how instances are placed on a physical server.

Dedicated instances are Amazon EC2 instances that run in a virtual private cloud (VPC) on hardware that's dedicated to a single customer. Your Dedicated Instances are physically isolated at the host hardware level from instances that belong to other AWS accounts. Dedicated Instances may share hardware with other instances from the same AWS account that are not Dedicated Instances

For more information, see:

- <http://aws.amazon.com/ec2/pricing/>
- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/instance-purchasing-options.html>
- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/dedicated-instance.html>

Networking Amazon VPC



26

Understand Virtual Private Cloud (VPC) concepts including:

- Subnets
- Security
- Networking
- VPN

Amazon Virtual Private Cloud (VPC)



Amazon
VPC

- Provision a **private, isolated virtual network** on the AWS cloud.
- Have complete control over your virtual networking environment.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and
Certification

27

With Amazon Virtual Private Cloud (VPC), you can define a virtual network topology that closely resembles a traditional network that you might operate in your own data center. You have complete control over your virtual networking environment, and you can easily customize the network configuration for your Amazon VPC, such as selection of IP address range, creation of subnets, configuration of route tables, and network gateways.

VPCs and Subnets



- A **subnet** defines a range of IP addresses in your VPC.
- You can launch AWS resources into a subnet that you select.
- A **private subnet** should be used for resources that won't be accessible over the Internet.
- A **public subnet** should be used for resources that will be accessed over the Internet.
- Each subnet must reside entirely within one Availability Zone and cannot span zones.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


 Training and
Certification

28

AWS assigns a unique ID to each subnet. Regardless of the type of subnet (public or private), the internal IP address range of the subnet is always private.

A public subnet has a route to an Internet gateway (i.e., for a web server accessible from the Internet).


A private subnet has no route to an Internet gateway (i.e., for a database server only accessed within the VPC).



Amazon Virtual Private Cloud, also known as Amazon VPC, allows you provision a logically isolated section of the AWS cloud where you can launch AWS resources in a virtual network that you define. You have complete control over your virtual networking environment, including selection of your own IP address range, creation of subnets, configuration of route tables, network access control lists, and network gateways. You can easily customize the network and configuration for your Amazon VPC instance. For example, you can create a public-facing subnet for your web servers that require access to the Internet and place your back-end systems, such as databases or application servers, in a private-facing subnet with no Internet access. You can leverage multiple layers of security, including security groups and network access control lists, to help control access to Amazon EC2 instances in each subnet. Additionally, you can create a hardware virtual private network (VPN) connection between your corporate data center and your VPC, which allows you to leverage the AWS cloud as an extension of your corporate data center.

Security in Your VPC

- Security groups
- Network access control lists (ACLs)
- Key Pairs



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

30

Amazon VPC provides various features that you can use to increase and monitor the security for your VPC:

- Security groups act as a firewall for associated Amazon EC2 instances, controlling both inbound and outbound traffic at the instance level.
- Network access controls lists (ACLs) act as a firewall for associated subnets, controlling both inbound and outbound traffic at the subnet level.
- Amazon EC2 uses public-key cryptography to encrypt and decrypt login information. Public-key cryptography uses a public key to encrypt a piece of data, and the recipient uses the private key to decrypt the data. The private and public keys are known as a *key pair*. To log in to your instance, you must create a key pair, specific the name of the key pair when you launch the instance, and provide the private key when you connect to the instance. Linux instances have no password, and you use a key pair to log in using SSH. Windows instances require a key pair to obtain the administrator password to log in using RDP.

VPN Connections



VPN Connectivity option	Description
AWS Hardware VPN	You can create an IPsec hardware VPN connection between your VPC and your remote network.
AWS Direct Connect	AWS Direct Connect provides a dedicated private connection from a remote network to your VPC.
AWS VPN CloudHub	You can create multiple AWS hardware VPN connections via your VPC to enable communications between various remote networks.
Software VPN	You can create a VPN connection to your remote network by using an Amazon EC2 instance in your VPC that's running a software VPN appliance .

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

31

You can connect your VPC to remote networks by using a VPN connection.

For more information, see:

- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpn-connections.html>
- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPN_CloudHub.html

Lab 1

Build your VPC and Launch a Web Server

45 minutes

Lab 1 Overview

- 1**
 - **Create a VPC**
 - Two public subnets
 - Two private subnets
 - Across two Availability Zones
- 2**
 - **Create an Application Server**
 - Create a security group for your instance.
 - Launch your instance.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



33


For more information, see:

AWS Web Hosting Best Practices: <https://aws.amazon.com/whitepapers/web-application-hosting-best-practices>

Hosting a Web App on Amazon Web Services:

<http://docs.aws.amazon.com/gettingstarted/latest/wah-linux/web-app-hosting-intro.html>

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



34

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



35

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

36

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



37

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



38

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



39

Lab 1 – Build your VPC and Launch a Web Server




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



40

Lab 1 – Build your VPC and Launch a Web Server



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



41

Storage Services

Amazon S3 and Amazon EBS



42

Understand storage options including:

- Amazon S3 (Requests, Buckets, Objects, Access, Protecting Data, Notifications, Replication, Request Routing, Optimization, Lifecycle Management with Glacier)
- Amazon EBS (Volumes, Snapshots, Optimization, Encryption, Performance)

Amazon Simple Storage Service (S3)



Amazon S3

- Storage for the Internet
- Natively online, HTTP access
- Storage that allows you to store and retrieve **any amount of data**, any time, from anywhere on the web
- **Highly scalable**, reliable, fast and durable

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



43

Amazon S3 is designed to make web-scale computing easier for developers. Amazon S3 provides a simple web services interface that can be used to store and retrieve any amount of data, at any time, from anywhere on the web. It gives any developer access to the same highly scalable, reliable, secure, fast, inexpensive infrastructure that Amazon uses to run its own global network of websites.

For more information, see: <http://aws.amazon.com/s3/>

Amazon S3 Facts



- Can store an **unlimited number of objects** in a bucket
- Objects can be **up to 5 TB**; no bucket size limit
- Designed for **99.999999999%** durability and **99.99%** availability of objects over a given year
- Can use **HTTP/S** endpoints to store and retrieve any amount of data, at any time, from anywhere on the web
- Is highly scalable, reliable, fast, and inexpensive
- Can use optional server-side **encryption** using AWS or customer-managed provided client-side encryption
- Auditing is provided by access logs
- Provides standards-based **REST** and **SOAP** interfaces

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



44

Here are some facts about Amazon S3. There is a 100-bucket limit per account. You can store an unlimited number of objects in a bucket. The size of an object can be up to 5 TB, and there is no limit to the size of a bucket. Amazon S3 is designed for 99.999999999% durability and 99.99% availability of objects over a given year. You can use HTTP or HTTPS endpoints to store and retrieve any amount of data, at any time, from anywhere on the web. Most importantly, Amazon S3 is highly scalable, reliable, fast, and inexpensive.

Common Use Scenarios



- Storage and backup
- Application file hosting
- Media hosting
- Software delivery
- Store AMIs and snapshots

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



45

Advanced use scenarios:

Using Amazon DevPay with Amazon S3: Amazon DevPay enables you to charge customers for using your Amazon S3 product through Amazon's authentication and billing infrastructure. You can charge any amount for your product, including usage charges (storage, transactions, and bandwidth), monthly fixed charges, and a one-time charge.


Publishing Content Using Amazon S3 and BitTorrent: You can direct your clients to your BitTorrent-accessible objects by giving them the .torrent file directly or by publishing a link to the BitTorrent URL of your object.

Hosting a Static Website on Amazon S3: You can host a static website on Amazon S3 by configuring a bucket for website hosting and then uploading your website content to the bucket.

For more information, see:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingDevPay.html>
- <http://docs.aws.amazon.com/AmazonS3/latest/dev/S3TorrentPublish.html>
- <http://docs.aws.amazon.com/AmazonS3/latest/dev/WebsiteHosting.html>

Amazon S3 Concepts



- Amazon S3 stores data as objects within **buckets**
- An object is composed of a file and optionally any **metadata** that describes that file
- You can have **up to 100 buckets** in each account
- You can **control access** to the bucket and its objects

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

46

To get the most out of Amazon S3, you need to understand a few simple concepts. First, Amazon S3 stores data as objects within buckets.

An object is composed of a file and any metadata that describes that file. To store an object in Amazon S3, you upload the file you want to store into a bucket. When you upload a file, you can set permissions on the object and add any metadata.

Buckets are logical containers for objects. You can have one or more buckets in your account. For each bucket, you can control access: in other words, who can create, delete and list objects in the bucket. You can also view access logs for the bucket and its objects and choose the geographical region where Amazon S3 will store the bucket and its contents.

Object Keys

An object key is the unique identifier for an object in a bucket.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



47

Because the combination of a bucket, key, and version ID uniquely identifies each object, Amazon S3 can be thought of as a basic data map between "bucket + key + version" and the object itself. Every object in Amazon S3 can be uniquely addressed through the combination of the web service endpoint, bucket name, key, and optionally, a version.

For example, in the URL <http://doc.s3.amazonaws.com/2006-03-01/AmazonS3.html>, "doc" is the name of the bucket and "2006-03-01/AmazonS3.html" is the key.

Amazon S3 Security



- You can **control access** to buckets and objects with:
 - Access Control Lists (ACLs)
 - Bucket policies
 - Identity and Access Management (IAM) policies
- You can upload or download data to Amazon S3 via **SSL** encrypted endpoints.
- You can **encrypt data** using AWS SDKs.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



48

Data Access:

- IAM policies: With IAM policies, you can only grant users within your own AWS account permission to access your Amazon S3 resources.
- ACLs: With ACLs, you can only grant other AWS accounts (not specific users) access to your Amazon S3 resources.
- Bucket Policies: Bucket policies in Amazon S3 can be used to add or deny permissions across some or all of the objects within a single bucket. Policies can be attached to users, groups, or Amazon S3 buckets, enabling centralized management of permissions. With bucket policies, you can grant users within your AWS account or another AWS account access to your Amazon S3 resources.

Data Transfer: For maximum security, you can securely upload data to and download data from Amazon S3 via the SSL encrypted endpoints. The encrypted endpoints are accessible both from the Internet and from within Amazon EC2 so that data is transferred securely both within AWS and to and from sources outside of AWS.

Data Storage: Amazon S3 provides multiple options for protecting data at rest. Customers who prefer to manage their own encryption keys can use a client encryption library such as the Amazon S3 Encryption Client to encrypt data before uploading to Amazon S3.

Alternatively, they can use Amazon S3 Server Side Encryption (SSE) if you prefer to have Amazon S3 manage encryption keys for you. With Amazon S3 SSE, you can encrypt data on upload simply by adding an additional request header when writing the object. Decryption happens automatically when data is retrieved.

Amazon S3 SSE uses one of the strongest block ciphers available: 256-bit Advanced Encryption Standard (AES-256). With Amazon S3 SSE, every protected object is encrypted with a unique encryption key. This object key itself is then encrypted with a regularly rotated master key. Amazon S3 SSE provides additional security by storing the encrypted data and encryption keys in different hosts. Amazon S3 SSE also makes it possible for you to enforce encryption requirements. For example, you can create and apply bucket policies that require that only encrypted data can be uploaded to your buckets.

Instead of using Amazon S3 SSE, you can also encrypt your data before sending it to Amazon S3. You can build your own library that encrypts your object data on the client side before uploading it to Amazon S3. Optionally, you can use an AWS SDK to automatically encrypt your data before uploading it to Amazon S3.

For more information, see:

<http://docs.aws.amazon.com/AmazonS3/latest/dev/UsingClientSideEncryption.html>

Amazon S3 Versioning

- Protects from **accidental overwrites and deletes** with no performance penalty.
- Generates a **new version with every upload**.
- Allows easily retrieval of deleted objects or **roll back** to previous versions.
- Three states of an Amazon S3 bucket
 - Un-versioned (default)
 - Versioning-enabled
 - Versioning-suspended



Versioning Enabled

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

49

Versioning is a method of keeping multiple variants of an object in the same bucket. You can use versioning to preserve, retrieve, and restore every version of every object stored in your Amazon S3 bucket. With versioning, you can easily recover from both unintended user actions and application failures.

In one bucket, for example, you can have two objects with the same key, but different version IDs, such as photo.gif (version 111111) and photo.gif (version 121212).

If you version-enable a bucket, it can never return to an unversioned state. You can, however, suspend versioning on that bucket.

Amazon S3 Object Lifecycle



Lifecycle management defines how Amazon S3 manages objects during their lifetime. Some objects that you store in an Amazon S3 bucket might have a well-defined lifecycle:

- Log files
- Archive documents
- Digital media archives
- Financial and healthcare records
- Raw genomics sequence data
- Long-term database backups
- Data that must be retained for regulatory compliance

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

50

Lifecycle management defines how Amazon S3 manages objects during their lifetime. Some objects that you store in an Amazon S3 bucket might have a well-defined lifecycle: if you are uploading periodic logs to your bucket, your application might need these logs for a week or a month after creation, and after that you might want to delete them. Some documents are frequently accessed for a limited period of time. After that, you might not need real-time access to these objects, but your organization might require you to archive them for a longer period and then optionally delete them.

Digital media archives, financial and healthcare records, raw genomics sequence data, long-term database backups, and data that must be retained for regulatory compliance are some of the kinds of objects that you might upload to Amazon S3 primarily for archival purposes.

When you configure a lifecycle rule, you specify the storage class you want to transition the object to and the number of days after object creation to transition it. You can transition objects to the Standard – Infrequent Access (IA) storage class, archive them to Amazon Glacier, or have them permanently deleted. Standard - IA is useful for data such as backups and other older, infrequently accessed data where high performance continues to be a requirement. It is suitable for objects greater than 128 kilobytes that you want to keep for at least 30 days. There is a retrieval fee associated with Standard - IA objects. For more information, see:

- [Amazon S3 Pricing](#)
- <http://aws.amazon.com/solutions/case-studies/yelp/>

Amazon S3 Pricing

- Pay only for what you use
- No minimum fee
- Prices based on location of your Amazon S3 bucket
- Estimate monthly bill using the **AWS Simple Monthly Calculator**
- Pricing is available as:
 - Storage Pricing
 - Request Pricing
 - Data Transfer Pricing: data transferred out of Amazon S3



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

51

Amazon S3 pricing is based on capacity and bandwidth actually used. Because Amazon S3 is an Internet-scale service that runs natively across an entire region, it can handle significant request throughput and bandwidth output. All bandwidth into Amazon S3 is free, but AWS charges a rate on bandwidth out. Most importantly, since Amazon S3 can handle any amount of data, it is important to note that you only pay for the amount of space you use. Prices are based on a prorated GB per month.

There is also a pricing calculator online as a reference. Note that pricing listed is in the US East (N. Virginia) Region at the time this training was developed.

For more information, see:

- Online Pricing Calculator: <http://calculator.s3.amazonaws.com/calc5.html>
- <https://aws.amazon.com/s3/pricing/>

Amazon Glacier



- Long term low-cost archiving service
- Optimal for infrequently accessed data
- Designed for 99.999999999% durability
- Three to five hours' retrieval time
- Less than \$0.01 per GB/month (depending on region)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

52

For more information about Glacier pricing, see:
<https://aws.amazon.com/glacier/pricing/>

Amazon S3 Storage Classes

Storage Class	Durability	Availability	Other Considerations
Amazon S3 Standard	99.999999999%	99.99%	
Amazon S3 Standard - Infrequent Access (IA)	99.999999999%	99.9%	<ul style="list-style-type: none">• Retrieval fee associated with objects• Most suitable for infrequently accessed data
Glacier	99.999999999%	99.99% (once restored)	<ul style="list-style-type: none">• Not available for real-time access• Must restore objects before you can access them• Restoring objects can take 3-5 hours

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



53

Each object in Amazon S3 has a storage class associated with it.

S3 Standard is ideal for performance-sensitive use cases and frequently used data. Standard is the default storage class in S3.

S3 Infrequent Access (IA) is optimized for long-lived and less frequently accessed data, such as backups and older data that are accessed less but still require high performance.

Glacier is suitable for archiving data where access is infrequent and a retrieval time of several hours is acceptable. Archived objects are not available for real-time access: they must be restored before they can be accessed. The Glacier storage class is very low-cost.

S3 Reduced Redundancy Storage (RSS) is designed for noncritical, reproducible data stored at lower levels of redundancy standards than the Standard or IA classes, thus reducing cost.

For more information, see:

- <http://docs.aws.amazon.com/AmazonS3/latest/dev/storage-class-intro.html>
- <https://aws.amazon.com/s3/storage-classes/>

Instructor Demo

Amazon S3



54

The instructor will demonstrate creating a bucket, setting access control for the bucket with a bucket policy, enabling versioning and server-side encryption, uploading an object, and setting the object lifecycle.

Amazon Elastic Block Store (EBS)



Amazon
EBS


- **Persistent block level storage** volumes offer consistent and low-latency performance.
- Stored data is automatically replicated within its Availability Zone.
- Snapshots are stored durably in Amazon S3.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



55

Amazon Elastic Block Store, also known as Amazon EBS, provides persistent block-level storage volumes for use with Amazon EC2 instances and offers consistent and low-latency performance. Amazon EBS is particularly suited for applications that require a database, file system, or access to raw block-level storage. Amazon EBS snapshots are durable and automatically replicated within their Availability Zone. Snapshots can be stored in Amazon S3.



Amazon EBS provides block-level storage volumes for use with Amazon EC2 instances. Amazon EBS volumes are highly available and reliable storage volumes that can be attached to any running instance in the same Availability Zone. The Amazon EBS volumes attached to an Amazon EC2 instance are exposed as storage volumes that persist independently from the life of the instance. When the volumes are not attached to an EC2 instance, you pay only for the cost of storage.

Amazon EBS Volume Types



- SSD-backed volumes are
 - Optimized for **transactional** workloads that involve **frequent read/write** operations with **small I/O size**.
 - Dominant in **IOPS** performance.
- HDD-backed volumes are
 - Optimized for **large streaming** workloads.
 - Dominant in **throughput** (measured in MiB/s).

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

57

Amazon EBS provides the following volume types, which differ in performance characteristics and price, so that you can tailor your storage performance and cost to the needs of your applications. The volume types fall into two categories: solid-state drives (SSD) and hard disk drives (HDD).

For more information, see:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

Amazon EBS Volume Types



	SSD		HDD	
Volume Type	General Purpose SSD (gp2)	Provisioned IOPS SSD (io1)	Throughput Optimized HDD (st1)	Cold HDD (sc1)
Description	Balances price and performance for a wide variety of transactional loads.	Highest-performance SSD volume designed for mission-critical applications.	Low-cost HDD designed for frequently accessed, throughput-intensive workloads.	Lowest cost HDD designed for less frequently accessed workloads.
Volume Sizes	1 GiB – 16 TiB	4 GiB – 16 TiB	500 GiB – 16 TiB	500 GiB – 16 TiB
Dominant Performance Attribute	IOPS	IOPS	MiB/s	MiB/s

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



58

For more information, see:

<https://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

Amazon EBS Facts



- EBS is recommended when data must be **quickly accessible** and requires **long-term persistence**.
- You can launch your EBS volumes as **encrypted** volumes – data stored at rest on the volume, disk I/O, and snapshots created from the volume are all encrypted.
- You can create **point-in-time snapshots** of EBS volumes, which are persisted to Amazon S3.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

59

You can mount multiple volumes on the same instance, but each volume can be attached to only one instance at a time. EBS volumes behave like raw, unformatted block devices. You can create a file system on top of these volumes or use them in any other way you would use a block device (like a hard drive).

Amazon EBS is recommended when data changes frequently and requires long-term persistence. EBS volumes are particularly well-suited for use as the primary storage for file systems and databases or for any applications that require fine granular updates and access to raw, unformatted, block-level storage. Amazon EBS is particularly helpful for database-style applications that frequently encounter many random reads and writes across the data set.

Amazon EBS Use Cases

- **OS:** Use for boot/root volume, secondary volumes
- **Databases:** Scales with your performance needs
- **Enterprise applications:** Provides reliable block storage to run mission-critical applications
- **Business continuity:** Minimize data loss and recovery time by regularly backing up using EBS Snapshots
- **Applications:** Install and persist any application



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

60

The Amazon EBS service is simply a virtual hard drive. So, a great use case for Amazon EBS is when you want the hard drive to persist past the life of the Amazon EC2 instance. Before Amazon EBS existed as a service, AWS only used physical local attached hard drives called *ephemeral storage*. The problem with that was that you couldn't stop an Amazon EC2 instance without losing all your data, because of the temporary nature of local storage.

That's why we created Amazon EBS to decouple the lifecycle of data persistence from the lifecycle of an EC2 instance. Amazon EBS volumes are ideal for root volumes you need to store and have block-level access to your operating system, database storage, and datasets that are smaller than 1 TB. Given its simple snapshot mechanism, Amazon EBS is a great use case for simplifying distributed backups as well.

For more information, see:

Dropcam Case Study using AWS and Amazon EBS -
<http://aws.amazon.com/solutions/case-studies/dropcam/>

Amazon EBS Pricing



Pay for what you provision:

- Pricing based on region
- Review Pricing Calculator online
- Pricing is available as:
 - Storage
 - IOPS

* Check Amazon EBS Pricing page for current pricing for all regions.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

61

Amazon EBS pricing is based on allocated storage, whether you use it or not. This is unlike Amazon S3, whose pricing is based on space actually in use. Prices may vary based on region or for IOPS.

For more information, see:

- Check online for current pricing for all regions -
<http://aws.amazon.com/ebs/pricing/>
- AWS Simple Monthly Calculator -
<http://calculator.s3.amazonaws.com/index.html>

Amazon EBS Scope



Amazon EBS volumes are in a single Availability Zone

EBS Volume 1



EBS Volume 2



Volume data is replicated across multiple servers in an Availability Zone.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

62

Amazon EBS volumes are designed to be highly available and reliable. Amazon EBS volume data is replicated across multiple servers in an Availability Zone to prevent the loss of data from the failure of any single component.

The durability of your volume depends on both the size of your volume and the percentage of the data that has changed since your last snapshot.

Amazon EBS volumes are designed for an annual failure rate (AFR) of between 0.1% and 0.2%, where failure refers to a complete or partial loss of the volume, depending on the size and performance of the volume. This is compared with commodity hard disks that will typically fail with an AFR of around 4%, making EBS volumes 10 times more reliable than typical commodity disks.

Because Amazon EBS servers are replicated within a single Availability Zone, mirroring data across multiple Amazon EBS volumes in the same Availability Zone will not significantly improve volume durability.

If you are interested in even more durability, with Amazon EBS you can create point-in-time consistent snapshots of your volumes that are then stored in Amazon S3 and automatically replicated across multiple Availability Zones.

Taking frequent snapshots of your volume is a convenient and cost-effective way to increase the long-term durability of your data. In the unlikely event that your Amazon EBS volume does fail, all snapshots of that volume will remain intact and will allow you to recreate your volume from the last snapshot point.

Amazon EBS and Amazon S3

	Amazon EBS 	Amazon S3 
Paradigm	Block storage with file system	Object store
Performance	Very fast	Fast
Redundancy	Across multiple servers in an Availability Zone	Across multiple facilities in a Region
Security	EBS Encryption – Data volumes and Snapshots	Encryption
Access from the Internet?	No (1)	Yes (2)
Typical use case	It is a disk drive	Online storage

- (1) Accessible from the Internet if mounted to server and set up as FTP, etc.
(2) Only with proper credentials, unless ACLs are world-readable

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



63

This table illustrates the significant differences between Amazon S3 and Amazon EBS. Amazon EBS volumes are network-attached hard drives that can be written to or read from at a block level. Amazon S3 is an object-level storage medium.

This means that you must write whole objects at a time. If you change one small part of a file, you must still rewrite the entire file in order to commit the change to Amazon S3. This can be very time-consuming if you have frequent writes to the same object.

Amazon S3 is optimized for write-once/read-many use cases. The other major difference is cost. With Amazon S3 you pay for what you use, and with Amazon EBS you pay for what you provision.

Amazon EC2 Instance Storage

- Is local, complimentary **direct attached block storage**.
- Includes availability, number of disks, and size **based on EC2 instance type**.
- Is optimized for **up to 365,000 Read IOPS** and 315,000 First Write IOPS.
- Is SSD or magnetic.
- Has **no persistence**.
- **Automatically deletes** data when an EC2 instance stops, fails or is terminated.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



64

An instance store provides temporary block-level storage for your instance. This storage is located on disks that are physically attached to the host computer. Instance store is ideal for temporary storage of information that changes frequently, such as buffers, caches, scratch data, and other temporary content, or for data that is replicated across a fleet of instances, such as a load-balanced pool of web servers.

For more information, see:

- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/InstanceStorage.html#instance-store-volumes>
- <http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/i2-instances.html>

Amazon EBS vs. Amazon EC2 Instance Store

Amazon EBS

- Data stored on an Amazon EBS volume can persist independently of the life of the instance.
- Storage is **persistent**.

Amazon EC2 Instance Store

- Data stored on a local instance store persists only as long as the instance is alive.
- Storage is **ephemeral**.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



65

Use the local instance store only for temporary data. For data that requires a higher level of durability, use Amazon EBS volumes or back up the data to Amazon S3. If you are using an Amazon EBS volume as a root partition and want your Amazon EBS volume to persist outside of the life of the instance, set the Delete on termination flag to “No”.

Reboot vs. Stop vs. Terminate

Characteristic	Reboot	Stop/Start (EBS-backed instances only)	Terminate
Host computer	The instance stays on the same host computer .	The instance runs on a new host computer .	
Public IP address	No change	New address assigned	
Elastic IP addresses (EIP)	EIP remains associated with the instance.	EIP remains associated with the instance.	EIP is disassociated from the instance.
Instance store volumes	Preserved	Erased	Erased
EBS volume	Preserved	Preserved	Boot volume is deleted by default .
Billing	Instance billing hour doesn't change.	You stop incurring charges as soon as state is changed to <i>stopping</i> .	You stop incurring charges as soon as state is changed to <i>shutting-down</i> .

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



66

The table shows the differences between rebooting, stopping, and terminating your instance.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/ec2-instance-lifecycle.html>

Knowledge Check

Q: What AWS service would help support your web application by **hosting static assets** and **storing user uploaded images and video off-instance**?

Amazon S3

Q: How would an EC2 instance find its private and public IP addresses?

Retrieve the instance metadata. <http://169.254.169.254/latest/meta-data/>

Q: What acts as an additional layer of security at the subnet level in a VPC?

Network ACLs

True or False: S3 limits the total amount you can store.

False (There is a 5TB limit per object)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




67

The answers to the knowledge check are shown in the slide.

Module 3

Security, Identity, and Access Management



When we talk about cloud security, we like to start with a discussion of the Shared Responsibility Model. Although AWS takes care of provisioning and maintaining the underlying cloud infrastructure, you will still need to perform several security configuration tasks to ensure that you stay safe in the cloud. The responsibility of AWS goes from the ground up to the hypervisor. AWS secures the hardware, software, facilities, and networks that run all products and services. Customers are responsible for securely configuring the services they sign up for and they put on those services.

AWS also performs the following responsibilities:

- Obtaining industry certifications and independent third-party attestations
- Publishing information about AWS security and control practices in whitepapers and web site content
- Providing certificates, reports, and other documentation directly to AWS customers under NDA (as required)

The amount of security configuration work you have to do varies, depending on how sensitive your data is and which services you select. For example, AWS services such as Amazon EC2 and Amazon S3 are completely under your control and require you to perform all of the necessary security configuration and management tasks. In the case of Amazon EC2, you are responsible for management of the guest OS (including updates and security patches), any application software or utilities you install on the instances, and the configuration of the AWS-provided firewall (called a security group) on each instance.

When you use any of the AWS managed services such as Amazon RDS, Amazon RedShift, or Amazon WorkDocs, you don't have to worry about launching and maintaining instances or patching the guest OS or applications—AWS handles that for you. For these managed services, basic security configuration tasks such as data backups, database replication, and firewall configuration happen automatically.

However, there are certain security features—such as IAM user accounts and credentials, SSL for data transmissions, and user activity logging—that you should configure no matter which AWS service you use.

AWS Support provides a highly personalized level of service for customers seeking technical help.

For more information, see:

- <https://aws.amazon.com/premiumsupport/>
- http://d0.awsstatic.com/whitepapers/Security/Intro_to_AWS_Security.pdf
- <https://d0.awsstatic.com/whitepapers/aws-security-whitepaper.pdf>

Physical Security

- 24/7 trained **security staff**
- AWS data centers in **nondescript** and **undisclosed** facilities
- **Two-factor authentication** for authorized staff
- **Authorization** for data center access



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
webservices | Training and
Certification

3

One of the main security responsibilities of AWS is the physical security of the data centers that house the AWS Cloud infrastructure. Amazon has many years of experience designing, constructing, and operating large-scale data centers.

The physical security measures that protect these data centers are some of the most comprehensive in the industry and include: 24/7 trained security guards; locations in nondescript, undisclosed facilities; two-factor authentication for ingress; authorization for data center access only for an approved, specific need; and continuous monitoring, logging, and auditing of physical access controls.

For more information, see: Security Center - <http://aws.amazon.com/security/>

Hardware, Software, and Network

- Automated **change-control** process
- Bastion servers that **record all access attempts**
- **Firewall** and other **boundary devices**
- AWS **monitoring** tools



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon web services | Training and Certification

4

The hardware and software that supports AWS Cloud services has been architected to be highly available, redundant, and extremely secure. All changes to AWS hardware and software are managed through a centralized, automated change control process, and all access to hardware or software must be authorized.

Privileged access to software and systems requires SSH logon and is allowed only through bastion servers that record all access attempts. AWS network devices, including firewall and other boundary devices, monitor and control communications at the external boundary of the network and at key internal boundaries.

AWS monitoring tools are designed to detect unusual or unauthorized activities and conditions at ingress and egress communication points. These tools monitor server and network usage, port scanning activities, application usage, and unauthorized intrusion attempts. AWS security monitoring tools help identify several types of denial of service (DoS) attacks, including distributed, flooding, and software/logic attacks.

For more information, see: <http://aws.amazon.com/security/>

Certifications and Accreditations

ISO 9001, ISO 27001, ISO 27017, ISO 27018, IRAP (Australia), MLPS Level 3 (China), MTCS Tier 3 Certification (Singapore) and more ...

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

5

AWS has successfully completed multiple audits, attestations, and certifications. AWS publishes a Service Organization Controls SOC 1 report, published under both the SSAE 16 and the ISAE 3402 professional standards, as SOC 2-Security and SOC 3 Report.

In addition, AWS has achieved ISO 9001, ISO 27001, ISO 27017, and ISO 27018 certifications, has been successfully validated as a Level 1 service provider under the Payment Card Industry (PCI) Data Security Standard (DSS), and currently offers HIPAA Business Associate Agreements to covered entities and their business associates subject to HIPAA.

In the realm of public sector certifications, AWS has achieved FedRAMP compliance, has received authorization from the U.S. General Services Administration to operate at the FISMA Moderate level, and is also the platform for applications with Authorities to Operate (ATOs) under the Defense Information Assurance Certification and Accreditation Program (DIACAP).

NIST, FIPS 140-2, CJIS, and DoD SRG Levels 2 and 4 are some of the other certifications AWS has received.

For more information, see: <http://aws.amazon.com/compliance/>

SSL Endpoints

SSL Endpoints	Security Groups	VPC
<p>Secure Transmission</p> <p>Use secure endpoints to establish secure communication sessions (HTTPS).</p>	<p>Instance Firewalls</p> <p>Use security groups to configure firewall rules for instances.</p>	<p>Network Control</p> <p>Use public and private subnets, NAT, and VPN support in your virtual private cloud to create low-level networking constraints for resource access.</p>

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



6

AWS provides customer access points, also called API endpoints, that allow HTTPS access so that you can establish secure communication sessions with your AWS services, including SSL and TLS. SSL encrypts the transmission, protecting each request or the response from being viewed in transit.

Security Groups

SSL Endpoints	Security Groups	VPC
Secure Transmission Use secure endpoints to establish secure communication sessions (HTTPS).	Instance Firewalls Use security groups to configure firewall rules for instances.	Network Control Use public and private subnets, NAT, and VPN support in your virtual private cloud to create low-level networking constraints for resource access.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.




7

AWS also provides security groups, which act like built-in firewalls for your virtual servers. You can control how accessible your instances are by configuring security group rules - from totally public to completely private, or somewhere in between. And when your instances reside within a virtual private cloud (VPC) subnet, you can control egress as well as ingress traffic.

Security groups can also be used by AWS services such as Amazon RDS, Amazon Redshift, Amazon EMR, and Amazon ElastiCache.

AWS Multi-Tier Security Groups



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification

8

You can set up security group rules for your EC2 instances to create a traditional multi-tiered web architecture:

The web tier security group can accept traffic on port 80/443 from anywhere on the Internet if you select source 0.0.0.0/0.

Similarly, the app tier can only accept traffic from the web tier, and the DB tier can only accept traffic from the app tier.

Lastly, we have also added a set of rules to allow remote administration over SSH port 22. We have restricted remote access by funneling all traffic through the app tier and allowing access only from a specific IP. After you use SSH to access an app tier server, you can then connect to machines on the web and DB security groups.

Amazon Virtual Private Cloud (VPC)

SSL Endpoints	Security Groups	VPC
Secure Transmission Use secure endpoints to establish secure communication sessions (HTTPS).	Instance Firewalls Use security groups to configure firewall rules for instances.	Network Control Use public and private subnets, NAT, and VPN support in your virtual private cloud to create low-level networking constraints for resource access.


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



9

The Amazon Virtual Private Cloud (VPC) service allows you to add another layer of network security to your instances by creating private subnets and even adding an IPsec VPN tunnel between your network and your VPC. Amazon VPC allows you to define your own network topology, including definitions for subnets, network access control lists, Internet gateways, routing tables, and virtual private gateways. The subnets that you create can be defined as either private or public.

For more information, see: <http://aws.amazon.com/vpc/>



Using AWS IAM, you can create and manage AWS users and groups and use permissions to allow and deny their access to AWS resources. You can use existing corporate identities to grant secure access to AWS resources, such as Amazon S3 buckets, without creating any new AWS identities.

For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_aws-services-that-work-with-iam.html

AWS IAM Authentication

- **Authentication**
- **AWS Management Console**
 - User Name and Password

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification 11

AWS services and resources can be accessed using the AWS Management Console, AWS CLI, or through SDKs and APIs from a wide range of supported platforms. Users and systems have to be authenticated before they can access AWS services and resources.

The AWS Management Console provides a web-based way to administer AWS services. If you're the account owner, you can sign in to the console directly using the Root Account. It is, however, advisable to create individual IAM users for each user and sign in using individual credentials.

IAM is a complimentary service.

For more information, see:

<http://docs.aws.amazon.com/IAM/latest/UserGuide/console.html>

AWS IAM Authentication

- **Authentication**
- **AWS CLI or SDK API**
 - Access Key and Secret Key

IAM User

Access Key ID: AKIAIOSFODNN7EXAMPLE
Secret Access Key: wJalrXUtnFEMI/K7MDENG/bPxRfCYEXAMPLEKEY

AWS CLI

```
:$ aws configure
AWS Access Key ID [*****O22A]:
AWS Secret Access Key [*****4m8i]:
Default region name [ap-southeast-1]:
Default output format [json]:
```

AWS SDK & API

Java Python .NET

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


amazon web services | Training and Certification

The AWS Command Line Interface is a unified tool to manage your AWS services. With AWS CLI, you can control multiple AWS services from the command line and automate them through scripts.

AWS CLI is supported on Windows, Linux, OS X, and Unix platforms.

AWS offers support for a wide variety of programming platforms, including .NET, Java, and Python.

For more information, see: <http://docs.aws.amazon.com/cli/latest/userguide/cli-chap-welcome.html>



As the number of users managing your AWS environment increases, it is helpful to manage permissions for multiple IAM users using IAM groups.

For more information, see:


For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_groups.html

AWS IAM Authorization

Authorization

- Policies:
 - Are JSON documents to describe permissions.
 - Are assigned to users, groups or roles.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon web services

Training and
Certification

14

After a user or system has been authenticated, they have to be authorized to access AWS services. To assign permissions to a user, group, role, or resource, you create a policy, which is a document that explicitly lists permissions.

An IAM role is similar to a user, in that it is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. However, instead of being uniquely associated with one person, a role is intended to be assumable by anyone who needs it. Also, a role does not have any credentials (password or access keys) associated with it. Instead, if a user is assigned to a role, access keys are created dynamically and provided to the user.

Policies and roles are covered in more detail in the following slides.

AWS IAM Policy Elements



```
{ "Version": "2012-10-17", "Statement": [ { "Sid": "Stmt1453690971587", "Action": [ "ec2:Describe*", "ec2:StartInstances", "ec2:StopInstances" ], "Effect": "Allow", "Resource": "*", "Condition": { "IpAddress": { "aws:SourceIp": "54.64.34.65/32" } } }, { "Sid": "Stmt1453690998327", "Action": [ "s3:GetObject*" ], "Effect": "Allow", "Resource": "arn:aws:s3:::example_bucket/*" } ] }
```

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



15

Policies are documents that are created using JavaScript Object Notation (JSON). A policy consists of one or more statements, each of which describes one set of permissions.

An IAM policy may consist of :

- **Version**
- **Id**
- **Statement**
- **Sid**
- **Effect**: Defines what the effect will be when the user requests access—either *allow* or *deny*. The default is that resources are denied to users, so you typically specify that you will allow users access to resource.
- **Principal**
- **NotPrincipal**
- **Actions**: Defines what actions you want to allow. Each AWS service has its own set of actions. Any actions that you do not explicitly allow are denied.
- **NotAction**
- **Resources**: Defines which resources you allow the action on. Users cannot access any resources that you have not explicitly granted permissions to.

- **NotResource**
- **Condition**
- **Supported Data Types**

AWS Policy Generator: You can use the AWS Policy Generator to easily generate policies.

AWS Policy Validator: The Policy Validator automatically examines your existing IAM access control policies to ensure that they comply with the IAM policy grammar.

AWS Policy Simulator: The simulator evaluates the policies that you choose and determines the effective permissions for each of the actions that you specify. The simulator uses the same policy evaluation engine that is used during real requests to AWS services.


Managed policies: Standalone policies that you can attach to multiple users, groups, and roles in your AWS account. Managed policies apply only to identities (users, groups, and roles) - not resources. You can use two types of managed policies:

- **AWS-managed policies:** Managed policies that are created and managed by AWS. If you are new to using policies, we recommend that you start by using AWS managed policies.
- **Customer-managed policies:** Managed policies that you create and manage in your AWS account. Using customer-managed policies, you have more precise control over your policies than when using AWS managed policies.


Inline policies: Policies that you create and manage, and that are embedded directly into a single user, group, or role.

For more information, see:

- AWS Policy Generator - <http://awspolicygen.s3.amazonaws.com/policygen.html>
- Access the Policy Simulator - <https://pollicysim.aws.amazon.com/home/index.jsp>
- Overview of IAM Policies -
http://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies.html
- IAM Policy Elements Reference -
http://docs.aws.amazon.com/IAM/latest/UserGuide/reference_policies_elements.html



IAM policies are assigned to IAM users and groups. These users are bound by the permissions defined in the IAM policy.



IAM policies may also be assigned to an IAM role.

For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html

AWS IAM Roles



- An IAM role uses a policy.
- An IAM role has no associated credentials.
- IAM users, applications, and services may assume IAM roles.



IAM Roles

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services

Training and
Certification


18

IAM policies may also be assigned to an IAM role.

An IAM role is similar to a user, in that it is an AWS identity with permission policies that determine what the identity can and cannot do in AWS. A role does not have any credentials (password or access keys) associated with it. Instead, if a user is assigned to a role, access keys are created dynamically and provided to the user.

For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles.html



You can use roles to delegate access to users, applications, or services that don't normally have access to your AWS resources.

Example: Application Access to AWS Resources



- Python application hosted on an Amazon EC2 Instance needs to interact with Amazon S3.
- AWS credentials are required:
 - Option 1: ~~Store AWS Credentials on the Amazon EC2 instance.~~
 - Option 2: Securely distribute AWS credentials to AWS Services and Applications.



IAM Roles

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


Amazon Web Services

20

In the example above, a custom application written in Python and hosted on an Amazon EC2 instance needs to interact with objects stored in an Amazon S3 bucket. Applications may access AWS resources in multiple ways. One way is to embed your AWS access key ID and secret access key in the application code or in a config file supported by the application. However, doing so may compromise the user's credentials. Changing or rotating the user's credentials would require an update in the code each time. This approach is not secure and feasible in many cases. The alternative and secure option is to use an IAM role to pass temporary security credentials as part of an instance profile.

For more information, see:

- Using Instance Profiles -
http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_use_switch-role-ec2_instance-profiles.html
- IAM Roles for Amazon EC2 -
<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html>



An instance profile is a container for an IAM role that you can use to pass role information to an EC2 instance when the instance starts.

In the example, an IAM role named PythonInEC2AccessS3 is created by an IAM user. The role grants access to an Amazon S3 bucket.


1. An application developer selects the PythonInEC2AccessS3 role while creating the Amazon EC2 instance. The instance will host a Python application that will need access to an Amazon S3 bucket.

Note An IAM role may be associated with an EC2 instance only during creation. The policy associated with the role can be modified at any time. A user launching an EC2 instance also needs appropriate permissions to associate an IAM role to the EC2 instance.

2. The Python application is installed on the EC2 instance. AWS SDK for Python (Boto3) is also installed on the instance. The application tries to access an Amazon S3 bucket. However, AWS credentials are not available on the instance.
3. The Python application uses the EC2 metadata service to gain access to temporary security credentials.
4. The application interacts with the Amazon S3 bucket specified in the *PythonInEC2AccessS3* role.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/iam-roles-for-amazon-ec2.html#instance-metadata-security-credentials>



IAM roles may also be associated with users.

In the above example, there are two AWS accounts A and B. *IAM User A-1* is part of *Account A* and *IAM User B-1* is part of *Account B*.

1. An IAM policy named *IAM Admin Policy* with access to an Amazon S3 bucket is associated to an IAM role named *IAM Admin Role*. User A-1 has an IAM policy with restricted access. This is done because *User A-1* does not normally need administrative privileges. However, User A-1 may sometimes have to perform tasks that require administrative privileges.
2. When required, *User A-1* assumes the *IAM Admin Role*. Doing so gives *User A-1* access to the S3 bucket. A user who assumes a role temporarily gives up his or her own permissions and instead takes on the permissions of the role. When the user exits, or stops using the role, the original user permissions are restored. It is therefore helpful to use IAM roles instead of changing the user's policies each time a change is required.

Note *User A-1*'s policy must contain permissions to assume the role.

3. *User A-1* gains access to the Amazon S3 bucket.
4. With IAM roles, you can establish trust relationships between your *trusting* account and other AWS *trusted* accounts. The trusting account owns the resource to be accessed and the trusted account contains the users who need access to the resource. *User B-1* from *Account B* assumes the *IAM Admin Role* from *Account A*.

5. *User B-1* gains access to the Amazon S3 bucket owned by *Account A*.

For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_create_for-user.html

Temporary Security Credentials (AWS STS)

The diagram shows a 'Session' box containing four items: 'Access Key ID' (key icon), 'Secret Access Key' (key icon), 'Session Token' (orange circle), and 'Expiration' (alarm clock icon). A bracket groups these four items under the label 'Temporary Security Credentials'. Another bracket indicates that the expiration time is '15 minutes to 36 hours'. To the right of the session box is a green dashed box containing a green key icon. Below the session box is a hand icon pointing at a screen with a downward arrow, representing a user interface.

Use Cases

- Cross account access
- Federation
- Mobile Users
- Key rotation for Amazon EC2-based apps

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


amazon web services | Training and Certification 23

AWS Security Token Service (AWS STS) provides trusted users with temporary security credentials that can control access to your AWS resources. These credentials are short-term and work almost identically to the long-term access key credentials. These credentials are generated dynamically and provided to the user when requested.

A session established with AWS STS consists of an access key ID, a secret access key, a session token, and an expiration time. The expiration time could be between 15 minutes and 36 hours. The keys are used to sign API requests and pass in the token as an additional parameter, which AWS uses to verify that the temporary access keys are valid.


For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_credentials_temp.html



AWS IAM is not appropriate for OS and application authentication.

AWS IAM Authentication and Authorization



The slide is titled "AWS IAM Authentication and Authorization". It features three sections: "Authentication" with two bullet points, "Authorization" with one bullet point, and a diagram of IAM entities. The "Authentication" section includes icons for "IAM User" (single person), "IAM Group" (two people), and "IAM Roles" (construction worker). A green key icon is in the top right corner, enclosed in a dashed box. The "Authorization" section includes a bullet point for "Policies". The footer contains the AWS logo and "Training and Certification" text, along with a copyright notice: "© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved." and page number "25".

Authentication

- **AWS Management Console**
 - User Name and Password
- **AWS CLI or SDK API**
 - Access Key and Secret Key

Authorization

- Policies

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

25

IAM is a powerful service to authenticate and authorize users and AWS resources.

AWS IAM Best Practices

- **Delete** AWS account (root) access keys.
- Create **individual** IAM users.
- **Use groups** to assign permissions to IAM users.
- Grant **least privilege**.
- Configure a **strong password policy**.
- Enable **MFA** for privileged users.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
web services

Training and
Certification

26

The slide shows some best practices to follow with IAM.

AWS IAM Best Practices (cont.)



- Use **roles for applications** that run on Amazon EC2 instances.
- Delegate by **using roles** instead of by sharing credentials.
- **Rotate credentials** regularly.
- **Remove unnecessary** users and credentials.
- Use **policy conditions** for extra security.
- **Monitor activity** in your AWS account.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon
Training and
Certification


27

For more information, see: <http://docs.aws.amazon.com/IAM/latest/UserGuide/best-practices.html>

AWS CloudTrail



- Records AWS API calls for accounts.
- Delivers log files with information to an Amazon S3 bucket.
- Makes calls using the AWS Management Console, AWS SDKs, AWS CLI and higher-level AWS services.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

28

AWS CloudTrail is a web service that records API calls to supported AWS services in your account and delivers log files to you.

Knowledge Check

Q: Your **web application** needs to **read/write** an Amazon DynamoDB table and an Amazon S3 bucket. This operation requires **AWS credentials** and **authorization to use AWS services**. What IAM entity should be used?

- User
- Group
- Role
- Policy

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



29

You have reached the end of this training module. Test out some of your new skills!

Instructor Demo

IAM



30

The instructor will demonstrate how to create a user, role, group, and policy using the IAM console.

Module 4: Databases



Understand the concepts of fundamental AWS database services including:

- Amazon Relational Database Service (RDS)
 - DB instances
 - Security groups
 - RDS Interfaces
- Amazon DynamoDB
 - DynamoDB data model
 - Supported operations
 - Provisioned throughput
 - Accessing DynamoDB

SQL and NoSQL Databases

	SQL	NoSQL
Data Storage	Rows and Columns	Key-Value
Schemas	Fixed	Dynamic
Querying	Using SQL	Focused on collection of documents
Scalability	Vertical	Horizontal

SQL

ISBN	Title	Author	Format
9182932465265	Cloud Computing Concepts	Wilson, Joe	Paperback
3142536475869	The Database Guru	Gomez, Maria	eBook

NoSQL

```
{  
  ISBN: 9182932465265,  
  Title: "Cloud Computing Concepts",  
  Author: "Wilson, Joe",  
  Format: "Paperback"  
}
```

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



2

A SQL database stores data in rows and columns. Rows contain all the information about one entry, and columns are the attributes that separate the data points. A SQL database schema is fixed: columns must be locked before data entry. Schemas can be amended if the database is altered entirely and taken offline. Data in SQL databases is queried using SQL (Structure Query Language), which can allow for complex queries. SQL databases scale vertically by increasing hardware power.

NoSQL databases store data using one of many storage models including key-value pairs, documents, and graphs. NoSQL schemas are dynamic, and information can be added rapidly. Each 'row' doesn't have to contain data for each 'column'. Data in NoSQL databases is queried by focusing on collections of documents. NoSQL databases scale horizontally, by increasing servers.

Data Storage Considerations


- No one size fits all.
- Analyze your data requirements by considering:
 - ✓ Data formats
 - ✓ Data size
 - ✓ Query frequency
 - ✓ Data access speed
 - ✓ Data retention period

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



3

No one size fits all when considering database types. You must take into consideration your data requirements, such as data formats, data size, the frequency of your queries, how quickly you need your data, and for how long you need to keep it.



This slide shows the database services AWS provides.

Amazon Relational Database Service (RDS)



Amazon
RDS

- Cost-efficient and **resizable capacity**
- Manages time-consuming **database administration** tasks
- Access to the full capabilities of **Amazon Aurora, MySQL, MariaDB, Microsoft SQL Server, Oracle, and PostgreSQL** databases

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



5

With Amazon RDS, you can access the full capabilities of a familiar MySQL, MariaDB, Microsoft SQL Server, Oracle, or PostgreSQL database. In addition, Amazon RDS for MySQL provides two distinct, but complementary, replication features: Multi-AZ deployments and read replicas that can be used in conjunction with each other to gain enhanced database availability, protect your latest database updates against unplanned outages, and scale beyond the capacity constraints of a single DB instance for read-heavy database workloads.

Amazon Aurora is a MySQL-compatible relational database engine that is part of Amazon RDS.

Amazon RDS

- Simple and **fast to deploy**
- Manages common database administrative tasks
- **Compatible** with your applications
- Fast, predictable performance
- Simple and **fast to scale**
- Secure
- Cost-effective



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

6

Amazon RDS is a web service that makes it easy to set up, operate, and scale a relational database in the cloud. It provides cost-efficient and resizable capacity while managing time-consuming database administration tasks, which allows you to focus on your applications and business. Amazon RDS gives you access to the full capabilities of a MySQL, Oracle, SQL Server, or Amazon Aurora database engine. This means that the code, applications, and tools you already use today with your existing databases can be used with Amazon RDS. Amazon RDS automatically patches the database software and backs up your database, storing the backups for a user-defined retention period and enabling point-in-time recovery. You benefit from the flexibility of being able to scale the compute resources or storage capacity associated with your relational database instance via a single API call.

DB Instances

- DB Instances are the basic building blocks of Amazon RDS.
- They are an **isolated database environment** in the cloud.
- They can **contain multiple user-created databases**.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

7

The basic building block of Amazon RDS is the DB instance. A DB instance is an isolated database environment in the cloud. A DB instance can contain multiple user-created databases, and you can access it by using the same tools and applications that you use with a stand-alone database instance. You can create and modify a DB instance by using the AWS Management Console, the Amazon AWS command line interface, or the Amazon RDS API.

How Amazon RDS Backups Work



Automatic Backups:

- Restore your database to a point in time.
- Are enabled by default.
- Let you choose a retention period up to 35 days.



Manual Snapshots:

- Let you build a new database instance from a snapshot.
- Are initiated by the user.
- Persist until the user deletes them.
- Are stored in Amazon S3.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

8


When automated backups are turned on for your DB instance, Amazon RDS automatically performs a full daily snapshot of your data (during your preferred backup window) and captures transaction logs (as updates to your DB instance are made). When you initiate a point-in-time recovery, transaction logs are applied to the most appropriate daily backup in order to restore your DB instance to the specific time you requested. Amazon RDS retains backups of a DB instance for a limited, user-specified period of time called the *retention period*, which by default is one day but can be set to up to thirty-five days.

Manual database snapshots are user-initiated and enable you to back up your DB instance in a known state as frequently as you want, and then restore to that specific state at any time. DB snapshots can be created with the AWS Management Console or CreateDBSnapshot API and are kept until you explicitly delete them with the Console or DeleteDBSnapshot API.

Manual database snapshots are kept in Amazon Simple Storage Service (Amazon S3). There is no additional charge for backup storage up to 100% of your consumed database storage for an active DB instance.

Cross-Region Snapshots

- Are a **copy** of a **database snapshot** stored in a **different AWS Region**.
- Provide a backup for disaster **recovery**.
- Can be used as a **base** for **migration** to a different region.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

9

Cross-region snapshot copy is available for all Amazon RDS engines. You can copy snapshots of any size. Copies can be moved between any of the public AWS Regions, and you can copy the same snapshot to multiple regions simultaneously by initiating more than one transfer. There is no charge for the copy operation itself; you pay only for the data transfer out of the source region and for the data storage in the destination region.

Amazon RDS Security



- Run your DB instance in an **Amazon VPC**.
- Use **IAM policies** to grant access to RDS resources.
- Use **Security Groups**.
- Use Secure Socket Layer (**SSL**) connections with DB instances (Amazon Aurora, Oracle, MySQL, MariaDB, PostgreSQL, Microsoft SQL Server).
- Use RDS **encryption** to secure instances and snapshots at rest.
- Use network encryption and transparent data encryption (**TDE**) with Oracle DB and Microsoft SQL Server instances.
- Use security features of your DB engine to **control access** to DB instance.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification

10

You can manage access to your Amazon Relational Database Service (Amazon RDS) resources and your databases on a DB instance. The method you use to manage access depends on what type of task the user needs to perform with Amazon RDS.


- Run your DB instance in an Amazon virtual private cloud (VPC) for the greatest possible network access control.
- Use AWS Identity and Access Management (IAM) policies to assign permissions that determine who is allowed to manage RDS resources. For example, you can use AWS IAM to determine who is allowed to create, describe, modify, and delete DB instances, tag resources, or modify DB security groups.
- Use security groups to control which IP addresses or EC2 instances can connect to your databases on a DB instance. When you first create a DB instance, its firewall prevents any database access except through rules specified by an associated security group.
- Use Secure Socket Layer (SSL) connections with DB instances running the MySQL, MariaDB, PostgreSQL, or Microsoft SQL Server database engines.

- Use Amazon RDS encryption to secure your RDS DB instances and snapshots at rest. Amazon RDS encryption uses the industry standard AES-256 encryption algorithm to encrypt your data on the server that hosts your RDS DB instance.
- Use network encryption and transparent data encryption with Oracle DB instances.
- Use the security features of your DB engine to control who can log in to the databases on a DB instance, just as you would if the database was on your local network.

For more information, see:

- Using Amazon RDS with Amazon Virtual Private Cloud (VPC) -
http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/USER_VPC.html
- Setting up an IAM user -
http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_SettingUp.html#CHAP_SettingUp.IAM
- Using SSL with a DB instance -
http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/UsingWithRDS_SSL.html
- Encrypting Amazon RDS Resources -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Overview.Encryption.html>
- Oracle NNE -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.html#Appendix.Oracle.Options.NetworkEncryption>
- Oracle TDE -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.html#Appendix.Oracle.Options.AdvSecurity>

A Simple Application Architecture



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

11

The slide shows a simple application stack with an application running in an Amazon EC2 instance supported by a master database running in an Amazon RDS database instance. Presenting the application behind an elastic load balancer allows for compute resiliency and scaling features such as Auto Scaling and ELB groups to be adopted in the future.

Multi-AZ RDS Deployment




- With **Multi-AZ** operation, your database is **synchronously replicated to another Availability Zone** in the same AWS Region.
- Failover** to the standby **automatically** occurs in case of master database failure.
- Planned maintenance is applied first to standby databases.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

12

Amazon RDS Multi-AZ deployments provide enhanced availability and durability for Database (DB) instances, making them a natural fit for production database workloads. When you provision a Multi-AZ DB instance, Amazon RDS automatically creates a primary DB instance and synchronously replicates the data to a standby instance in a different Availability Zone (AZ). Each Availability Zone runs on its own physically distinct, independent infrastructure and is engineered to be highly reliable. In case of an infrastructure failure (for example, instance hardware failure, storage failure, or network disruption), Amazon RDS performs an automatic failover to the standby so that you can resume database operations as soon as the failover is complete. Because the endpoint for your DB instance remains the same after a failover, your application can resume database operation without the need for manual administrative intervention.



The slide shows an application stack that uses AWS reliability and durability features. An ELB group of Amazon EC2 instances supports the application logic. The instances use a Multi-AZ Amazon RDS deployment. In the event of infrastructure failure, the database fails over to a standby instance. The application logic retries its database connections to the same endpoint as before, and the service resumes using the new master. Meanwhile, a new standby is instantiated.

In addition to Amazon RDS's automatic backups, the database snapshot feature is used to ensure that backups are durably retained. You can create a new database instance from a database snapshot whenever you want.

For more information, see:

<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Concepts.MultiAZ.html>

Amazon RDS Best Practices



- **Monitor** your memory, CPU, and storage usage.
- Use **Multi-AZ** deployments to automatically provision and maintain a synchronous standby in a different Availability Zone.
- Enable **automatic backups**.
- Set the **backup window** to occur during the daily low in WriteIOPS.
- To increase the I/O capacity of a DB instance:
 - Migrate to a DB instance class with high I/O capacity.
 - Convert from standard storage to provisioned IOPS storage and use a DB instance class optimized for **provisioned IOPS**.
 - Provision additional throughput capacity (if using provisioned IOPS storage).
- If your client application is caching the DNS data of your DB instances, set a TTL of less than 30 seconds.
- **Test** failover for your DB instance.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



14

- Monitor your memory, CPU, and storage usage. Amazon CloudWatch can be set up to notify you when usage patterns change or when you approach the capacity of your deployment, so that you can maintain system performance and availability.
- Use Multi-AZ deployments to automatically provision and maintain a synchronous standby replica in a different Availability Zone.
- Enable automatic backups, and set the backup window to occur during the daily low in WriteIOPS.
- On a MySQL DB instance:
 - Do not create more than 10,000 tables using provisioned IOPS (input/output operations per second) or 1000 tables using standard storage. Large numbers of tables will significantly increase database recovery time after a failover or database crash. If you need to create more tables than recommended, set the `innodb_file_per_table` parameter to 0.
 - Avoid tables in your database growing too large. Underlying file system constraints restrict the maximum size of a MySQL table file to 2 TB.

Instead, partition your large tables so that file sizes are well under the 2 TB limit. This approach can also improve performance and recovery time.

- If your database workload requires more I/O than you have provisioned, recovery after a failover or database failure will be slow. To increase the I/O capacity of a DB instance, do any or all of the following:
 - Migrate to a DB instance class with high I/O capacity.
 - Convert from standard storage to provisioned IOPS storage, and use a DB instance class that is optimized for provisioned IOPS.
- If you are already using provisioned IOPS storage, provision additional throughput capacity.
- If your client application is caching the DNS data of your DB instances, set a time to live (TTL) of less than 30 seconds. Because the underlying IP address of a DB instance can change after a failover, caching the DNS data for an extended time can lead to connection failures if your application tries to connect to an IP address that is no longer in service.
- Test failover for your DB instance to understand how long the process takes for your use case and to ensure that the application that accesses your DB instance can automatically connect to the new DB instance after failover.

For more information, see:

http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/CHAP_BestPractices.html

Amazon DynamoDB



Amazon
DynamoDB

- Allows you to store any amount of data with **no limits**.
- Provides fast, predictable performance using **SSDs**.
- Allows you to easily provision and change the **request capacity** needed for each table.
- Is a **fully managed, NoSQL** database service.


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Training and
Certification

15

Amazon DynamoDB is a fully managed NoSQL database service that offers high performance, predictable throughput, and low cost. It is easy to set up, operate, and scale. With Amazon DynamoDB, you can start small, specify the throughput and storage you need, and easily scale your capacity requirements in seconds, as needed. It automatically partitions data over multiple servers to meet your requested capacity. In addition, Amazon DynamoDB automatically replicates your data synchronously across multiple Availability Zones within an AWS Region to ensure high availability and data durability.



In Amazon DynamoDB, a *table* is a collection of *items*, and each item is a collection of *attributes*. Each attribute in an item is a name-value pair. An attribute can be a scalar (single-valued), a JSON document, or a set.

For more information, see:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.html>

Primary Keys

The diagram illustrates a table structure for 'Table: Music'. The table has five columns: 'Artist' (Partition Key), 'Song Title' (Sort Key), 'Album Title', 'Year', and 'Genre'. The first column, 'Artist', is highlighted in orange and labeled 'Partition Key'. The second column, 'Song Title', is highlighted in blue and labeled 'Sort Key'. A bracket on the left side of the table is labeled 'Table: Music'. To the right of the table, text specifies 'Table: Music', 'Partition Key: Artist', and 'Sort Key: Song Title'. Above the table, there is a small icon of a blue database server.

(DynamoDB maintains a sorted index for both keys)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon | Training and Certification

17

When you create a table, in addition to the table name, you must specify the primary key of the table. As in other databases, a primary key in DynamoDB uniquely identifies each item in the table, so that no two items can have the same key. When you add, update, or delete an item in the table, you must specify the primary key attribute values for that item.

DynamoDB supports two different kinds of primary keys:

1. Partition Key: A simple primary key, composed of one attribute known as the partition key. DynamoDB uses the partition key's value as input to an internal hash function; the output from the hash function determines the partition where the item is stored. No two items in a table can have the same partition key value.
2. Partition Key and Sort Key: A composite primary key, composed of two attributes. The first attribute is the partition key, and the second attribute is the sort key. DynamoDB uses the partition key value as input to an internal hash function; the output from the hash function determines the partition where the item is stored. All items with the same partition key are stored together, in sorted order by sort key value. It is possible for two items to have the same partition key value, but those two items must have different sort key values.

For more information, see:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/HowItWorks.html>

Provisioned Throughput



- You specify how much **provisioned throughput capacity** you need for reads and writes.
- Amazon DynamoDB allocates the necessary machine resources to meet your needs.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

18

When you create or update a table, you specify how much provisioned throughput capacity you need for reads and writes. Amazon DynamoDB will allocate the necessary machine resources to meet your throughput needs while ensuring consistent, low-latency performance.

A unit of *read capacity* represents one strongly consistent read per second (or two eventually consistent reads per second) for items as large as 4 KB. A unit of *write capacity* represents one write per second for items as large as 1 KB.

Supported Operations



- **Query:**
 - Query a table using the partition key and an optional sort key filter.
 - If the table has a secondary index, query using its key.
 - It is the **most efficient way to retrieve items** from a table or secondary index.
- **Scan:**
 - You can scan a table or secondary index.
 - Scan reads every item – **slower than querying**.
- You can use conditional expressions in both Query and Scan operations.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



19


The Query operation enables you to query a table using the partition key and an optional sort key filter. If the table has a secondary index, you can also query the index using its key. You can query only tables that have a composite primary key (partition key and sort key). You can also query any secondary index on such tables. Query is the most efficient way to retrieve items from a table or a secondary index.

Amazon DynamoDB also supports a Scan operation, which you can use on a table or a secondary index. The Scan operation reads every item in the table or secondary index. For large tables and secondary indexes, a scan can consume a large amount of resources; for this reason, we recommend that you design your applications so that you can use the Query operation mostly, and use Scan only where appropriate.

You can use conditional expressions in both the Query and Scan operations to control which items are returned.

For more information, see:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/QueryAndScan.html>



The slide shows a simple application architecture using Amazon DynamoDB to store data processed by applications on Amazon EC2 instances.

Amazon RDS and Amazon DynamoDB

Factors	Relational (Amazon RDS)	NoSQL (Amazon DynamoDB)
Application Type	<ul style="list-style-type: none">Existing database appsBusiness process-centric apps	<ul style="list-style-type: none">New web-scale applicationsLarge number of small writes and reads
Application Characteristics	<ul style="list-style-type: none">Relational data models, transactionsComplex queries, joins, and updates	<ul style="list-style-type: none">Simple data models, transactionsRange queries, simple updates
Scaling	Application or DBA -architected (clustering, partitions, sharding)	Seamless, on-demand scaling based on application requirements
QoS	<ul style="list-style-type: none">Performance—depends on data model, indexing, query, and storage optimizationReliability and availabilityDurability	<ul style="list-style-type: none">Performance—Automatically optimized by the systemReliability and availabilityDurability

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



21

One type does not fit all. The choice depends on several factors. You can use both relational and NoSQL databases in one application, depending on requirements. This table provides a side-by-side comparison of relational and non-relational databases.

Database Considerations

If You Need	Consider Using
A relational database service with minimal administration	Amazon RDS <ul style="list-style-type: none">Choice of Amazon Aurora, MySQL, MariaDB, Microsoft SQL Server, Oracle, or PostgreSQL database enginesScale compute and storageMulti-AZ availability 
A fast, highly scalable NoSQL database service	Amazon DynamoDB <ul style="list-style-type: none">Extremely fast performanceSeamless scalability and reliabilityLow cost 
A database you can manage on your own	Your choice of AMIs on Amazon EC2 and Amazon EBS that provide scale compute and storage, complete control over instances, and more. 

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



22

AWS provides several database alternatives for developers. You can run fully managed relational and NoSQL services, or you can operate your own database in the cloud on Amazon EC2 and Amazon EBS. If you need a relational database service with minimal administration, consider using Amazon RDS. If you need a fast, highly scalable NoSQL database service, consider using Amazon DynamoDB. If you need a relational database you can manage on your own, consider using your choice of relational AMIs.

Knowledge Check

Q: What are the basic building blocks of Amazon Relational Database Service (RDS)?

DB Instances

True or False: Amazon DynamoDB allows you to store any amount of data with no limits.

True

True or False: Scan is the most efficient way to retrieve items from a DynamoDB table.

False

Q: You are creating a resilient, durable application using Amazon RDS. In addition to Amazon RDS's automatic backups, what feature should you use to ensure that your backups are durable retained?

Manual Snapshots

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Training and
Certification

23


The answers to the knowledge check questions are shown on the slide.

Lab 2

**Build your database server and
interact with your database
using an application.**

45 minutes

What You're Starting With – Lab 1 Review



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



25

Lab 2 Overview

1

- Create a **database server**
- Create a **security group**
- Create a **DB subnet group**
- Create an **Amazon RDS DB instance**
- Get database connection string

While waiting for the database to start, grab a coffee!

2

- Open a web application from a browser
- Insert DB connection string
- App will **populate a table** with records
- App will **display records** for a table


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

26

Lab 2 – Build Your Database Server and Connect to It

Your lab environment starts with the following resources




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



27

Lab 2 – Build Your Database Server and Connect to It

You will create a security group for the RDS instances




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

28

Lab 2 – Build Your Database Server and Connect to It




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



29

Lab 2 – Build Your Database Server and Connect to It




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

30

Lab 2 – Build Your Database Server and Connect to It



~ 45 min.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

31


Module 5

AWS Elasticity and Management Tools



This module describes the fundamental elements of AWS Management Tools.

Triad of Services



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

2

Auto Scaling works as a triad of services working in sync. Elastic Load Balancing and EC2 instances feed metrics to Amazon CloudWatch. Auto Scaling defines a group with launch configurations and Auto Scaling policies. Amazon CloudWatch alarms execute Auto Scaling policies to affect the size of your fleet. All of these services work well individually, but together they become more powerful and increase the control and flexibility our customers demand.

Elastic Load Balancing



Elastic Load
Balancing

- **Distributes** traffic across multiple EC2 instances, in multiple Availability Zones
- Supports **health checks** to detect unhealthy Amazon EC2 instances
- Supports the **routing and load balancing** of HTTP, HTTPS, SSL, and TCP traffic to Amazon EC2 instances

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



3

Understand Elastic Load Balancing (ELB) concepts including:

- Classic Load Balancer
- Application Load Balancer
- Load Balancer Comparison

Elastic Load Balancing automatically distributes incoming application traffic across multiple Amazon EC2 instances. It enables you to achieve greater fault tolerance in your applications, and seamlessly provides the amount of load balancing capacity needed in response to incoming application traffic. Elastic Load Balancing detects unhealthy instances within a pool and automatically reroutes traffic to healthy instances until the unhealthy instances have been restored. You can enable Elastic Load Balancing within a single Availability Zone or across multiple zones for even more consistent application performance.

You can access and work with your load balancer using one of the following interfaces:

- AWS Management Console: A simple web browser interface that you can use to create and manage your load balancers without using additional software or tools.


- Command Line Interfaces: A Java-based command line client that wraps the SOAP API.
- AWS SDKs: Language-specific APIs that take care of many of the connection details, such as calculating signatures, handling request retries, and error handling.
- Query API: Low-level API actions that you call using HTTPS requests.

For more information, see:

- AWS SDKs: <https://aws.amazon.com/tools/#SDKs>
- Application Load Balancer Query API:
<http://docs.aws.amazon.com/elasticloadbalancing/latest/APIReference/Welcome.html>
- Classic Load Balancer Query API:
<http://docs.aws.amazon.com/elasticloadbalancing/2012-06-01/APIReference/Welcome.html>

Classic Load Balancer - How It Works

Register instances with your load balancer.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

4

A load balancer distributes incoming application traffic across multiple EC2 instances in multiple Availability Zones. This increases the fault tolerance of your applications. Elastic Load Balancing detects unhealthy instances and routes traffic only to healthy instances.

Your load balancer serves as a single point of contact for clients. This increases the availability of your application. You can add and remove instances from your load balancer as your needs change, without disrupting the overall flow of requests to your application. Elastic Load Balancing scales your load balancer as traffic to your application changes over time. Elastic Load Balancing can scale to the vast majority of workloads automatically.


A listener checks for connection requests from clients, using the protocol and port that you configure, and forwards requests to one or more registered instances using the protocol and port number that you configure. You add one or more listeners to your load balancer.

You can configure health checks, which are used to monitor the health of the registered instances, so that the load balancer can send requests only to the healthy instances.

By default, the load balancer distributes traffic evenly across the Availability Zones that you enable for your load balancer. To distribute traffic evenly across all registered instances in all enabled Availability Zones, enable *cross-zone load balancing* on your load balancer.

For more information, see:

<http://docs.aws.amazon.com/elasticloadbalancing/latest/classic/introduction.html>



An Application Load Balancer (ALB) functions at the application layer (layer 7 of the OSI model). The ALB makes routing decisions based on the content of the application traffic in the HTTP messages.

With ALB, we require you to enable multiple Availability Zones.

A load balancer serves as the single point of contact for clients. You add one or more listeners to your load balancer.

A listener checks for connection requests from clients, using the protocol and port that you configure, and forwards requests to one or more target groups, based on the rules that you define. Each rule specifies a target group, condition, and priority. When the condition is met, the traffic is forwarded to the target group. You must define a default rule for each listener, and you can add rules that specify different target groups based on the content of the request (also known as *content-based routing*).

Each target group routes requests to one or more registered targets, such as EC2 instances, using the protocol and port number that you specify. You can register a target with multiple target groups. You can configure health checks on a per target group basis. Health checks are performed on all targets registered to a target group that is specified in a listener rule for your load balancer.

The slides shows the basic components of an ALB. Each listener contains a default rule, and one listener contains another rule that routes requests to a different target group. One target is registered with two target groups. The listener on the right has been configured with a rule to route /api requests to one target and /mobile requests to another.

For more information, see:

<http://docs.aws.amazon.com/elasticloadbalancing/latest/application/introduction.html>

Load Balancer Comparison



Classic Load Balancer

benefits include support for:

- EC2-Classic.
- VPC.
- TCP and SSL listeners.
- Sticky sessions.
- OSI Layer 4
(network protocol level)

ALB benefits include support

for:

- Path-based routing.
- Routing requests to multiple services on a single EC2 instance.
- Containerized applications.
- Monitoring the health of each service independently.
- OSI Layer 7
(application level)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

6

ALB access logs contain additional information and are stored in compressed format, and ALB has improved load balancer performance.

CLB is an OSI Layer 4 load balancer that works at the network protocol level and does not look inside of the actual network packets, remaining unaware of the specifics of HTTP and HTTPS. In other words, it balances the load without necessarily knowing a whole lot about it.

ALB is an OSI Layer 7 load balancer, which is more sophisticated and more powerful. It inspects packets, has access to HTTP and HTTPS headers, and (armed with more information) can do a more intelligent job of spreading the load out to the target.

For more information, see:

<http://docs.aws.amazon.com/elasticloadbalancing/latest/userguide/what-is-load-balancing.html>

Amazon CloudWatch



Amazon
CloudWatch

- A **monitoring service** for AWS cloud resources and the applications you run on AWS
- **Visibility into** resource utilization, operational performance, and overall demand patterns
- **Custom application-specific** metrics of your own
- **Accessible** via AWS Management Console, APIs, SDK, or CLI

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification


7

Amazon CloudWatch lets you view graphs, set alarms to troubleshoot, spot trends, and take automated action based on the state. It is accessible via the AWS Management Console, APIs, SDK, or CLI. You can customize with your own metrics or use a sample template found online.

Amazon CloudWatch Facts



- Monitor other AWS resources
 - View graphics and statistics
- Set Alarms



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification


8

For Amazon EC2 instances, Amazon CloudWatch basic monitoring collects and reports metrics for CPU utilization, data transfer, and disk usage activity from each Amazon EC2 instance at a five-minute frequency. Amazon CloudWatch detailed monitoring provides these same metrics at one-minute intervals, and also enables data aggregation by Amazon EC2 AMI ID and instance type.

Set alarms on any of your metrics to receive notifications. You can also use Auto Scaling to add or remove Amazon instances.

For more information, see:

- <http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/AlarmThatSendsEmail.html>
- <https://aws.amazon.com/cloudwatch/details/#other-aws-resource-monitoring>
- <https://aws.amazon.com/blogs/aws/new-cloudwatch-events-track-and-respond-to-changes-to-your-aws-resources/>



Amazon CloudWatch is a metrics repository. AWS products put metrics into the repository, and you retrieve statistics based on the metrics. Statistics can be graphically presented in the CloudWatch console.

CloudWatch Metrics Examples

CloudWatch Metrics by Category

Your CloudWatch metric summary has loaded. Total metrics: 97

EBS Metrics: 24 EC2 Metrics: 38 S3 Metrics: 16 SNS Metrics: 3 SQS Metrics: 14
Per-Volume Metrics: 24 Per-Instance Metrics: 38 Storage Metrics: 18 Topic Metrics: 3 Queue Metrics: 14

EC2 > Per-Instance Metrics

Showing all results (100) for cpu. For more results expand your search to All EC2 Metrics.
Select All | Clear

InstanceId	InstanceName	Metric Name
i-2be9e804		CPUUtilization
i-2c7e7653		CPUUtilization
i-5b210d23		CPUUtilization
i-8da54f60		CPUUtilization


Actions: Add to Dashboard, Copy URL, Create Alarm

Time Range: Relative, Absolute, UTC (GMT)
From: 12 hours ago To: 0 hours ago
Zoom: 1h | 3h | 6h | 12h | 1d | 3d | 1w | 2w

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

10



The slide shows screenshots from the Amazon CloudWatch console. In the example, the user has selected an EC2 per-instance metric of CPUUtilization.

Auto Scaling



Auto
Scaling

- **Scale** your Amazon EC2 capacity **automatically**
- Well-suited for applications that experience **variability in usage**
- Available at no additional charge

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification


11

Understand Auto Scaling concepts including:

- Launch configurations
- Auto Scaling groups
- Scaling plans
- Auto Scaling lifecycle
- Auto Scaling limits

Auto Scaling helps you ensure that you have the correct number of EC2 instances available to handle the load for your application. Auto Scaling is particularly well-suited for applications that experience hourly, daily, or weekly variability in usage.

Auto Scaling Benefits



Better Fault Tolerance



Better Availability



Better Cost Management



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

12

Adding Auto Scaling to your application architecture is one way to maximize the benefits of the AWS Cloud. When you use Auto Scaling, your applications gain the following benefits:

- Better fault tolerance: Auto Scaling can detect when an instance is unhealthy, terminate it, and launch an instance to replace it. You can also configure Auto Scaling to use multiple Availability Zones. If one Availability Zone becomes unavailable, Auto Scaling can launch instances in another one to compensate.
- Better availability: Auto Scaling can help you ensure that your application always has the right amount of capacity to handle the current traffic demands.
- Better cost management: Auto Scaling can dynamically increase and decrease capacity as needed. Because you pay for the EC2 instances you use, you save money by launching instances when they are actually needed and terminating them when they aren't needed.

Launch Configurations

- A **launch configuration** is a template that an Auto Scaling group uses to launch EC2 instances.
- When you create a launch configuration, you can specify:
 - AMI ID
 - Instance type
 - Key pair
 - Security groups
 - Block device mapping
 - User data



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

13

When you create an Auto Scaling group, you must specify a launch configuration. You can specify your launch configuration with multiple Auto Scaling groups. However, you can only specify one launch configuration for an Auto Scaling group at a time, and you can't modify a launch configuration after you've created it. If you want to change the launch configuration for your Auto Scaling group, you must create a new launch configuration and then update your Auto Scaling group with the new launch configuration. When you change the launch configuration for your Auto Scaling group, any new instances are launched using the new configuration parameters, but existing instances are not affected.


For more information, see:

<http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/LaunchConfiguration.html>

Auto Scaling Groups



- Contain a collection of EC2 instances that share similar characteristics.
- Instances in an Auto Scaling group are treated as a **logical grouping** for the purpose of instance scaling and management.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

14

You can create collections of EC2 instances, called Auto Scaling groups. You can specify the minimum number of instances in each Auto Scaling group, and Auto Scaling ensures that your group never goes below this size. You can specify the maximum number of instances in each Auto Scaling group, and Auto Scaling ensures that your group never goes above this size. If you specify the desired capacity, either when you create the group or at any time thereafter, Auto Scaling ensures that your group has this many instances. If you specify scaling policies, Auto Scaling can launch or terminate instances as demand on your application increases or decreases.

Dynamic Scaling



- You can create a scaling policy that uses **CloudWatch alarms** to determine:
 - When your Auto Scaling group should **scale out**.
 - When your Auto Scaling group should **scale in**.
- You can use alarms to monitor:
 - Any of the metrics that AWS services send to Amazon CloudWatch.
 - Your own **custom metrics**.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


 Training and Certification

15

Each CloudWatch alarm watches a single metric and sends messages to Auto Scaling when the metric breaches a threshold that you specify in your policy.

For more information, see:

- <http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/as-scale-based-on-demand.html>
- http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/policy_creating.html



The slide shows the basic lifecycle of instances within an Auto Scaling group.

1. The Scaling group has a desired capacity of three instances.
2. A CloudWatch alarm triggers scaling events, and policies scale the group at specific dates and times.
3. The scaling policy launches an instance and attaches it to the Auto Scaling group.
4. A health check fails and triggers an alarm similar to scaling out.
5. The instance is terminated.
6. The instance is detached from the Auto Scaling group.

AWS Trusted Advisor



AWS Trusted Advisor

- **Best practice** and recommendation engine.
- Provides AWS customers with performance and security recommendations in four categories:
 - Cost optimization
 - Security
 - Fault tolerance
 - Performance improvement.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

17

The status of checks provided by AWS Trusted Advisor is shown by using color coding on the dashboard page:

- Red: action recommended
- Yellow: investigation recommended
- Green: no problem detected

For each check, you can review a detailed description of the recommended best practice, a set of alert criteria, guidelines for action, and a list of useful resources on the topic.

Cost Optimization

- Amazon EC2 Reserved Instance Optimization
- Low-utilization Amazon EC2 Instances
- Idle load balancers
- Underutilized Amazon EBS volumes
- Unassociated Elastic IP addresses
- Amazon RDS idle DB instances

Cost Optimization



2 ✓ 4 ▲
0 !
0 excluded items

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification 18

AWS Trusted Advisor helps you save money on AWS by checking for unused and idle resources and making commitments to reserved capacity.

The following cost optimization checks are available with Trusted Advisor:

- Amazon EC2 Reserved Instance Optimization: Checks your Amazon Elastic Compute Cloud (Amazon EC2) computing consumption history and calculates an optimal number of Partial Upfront Reserved Instances. Recommendations are based on the previous calendar month's hour-by-hour usage aggregated across all consolidated billing accounts.
- Low Utilization Amazon EC2 Instances: Checks the Amazon EC2 instances that were running at any time during the last 14 days and alerts you if the daily CPU utilization was 10% or less and network I/O was 5 MB or less on four or more days.
- Idle Load Balancers: Checks your Elastic Load Balancing configuration for load balancers that are not actively used.
- Underutilized Amazon EBS Volumes: Checks Amazon Elastic Block Store (Amazon EBS) volume configurations and warns when volumes appear to be underused.

- **Unassociated Elastic IP Addresses:** Checks for Elastic IP addresses (EIPs) that are not associated with a running Amazon EC2 instance.
- **Amazon RDS Idle DB Instances:** Checks the configuration of your Amazon Relational Database Service (Amazon RDS) for any DB instances that appear to be idle. If a DB instance has not had a connection for a prolonged period of time, you can shut down the instance to reduce costs. If persistent storage is needed for data on the instance, you can use lower-cost options such as taking and retaining a DB snapshot.

Security

- Security groups
- AWS IAM use
- Amazon S3 bucket permissions
- MFA on Root Account
- AWS IAM password policy
- Amazon RDS security group access risk



Security



4 ✓ 2 ▲

3 !

1 excluded items



19

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

AWS Trusted Advisor helps you improve the security of your application by closing gaps, enabling various AWS security features, and examining your permissions.

The following security checks are available with Trusted Advisor:

- Security Groups - Specific Ports Unrestricted: Checks security groups for rules that allow unrestricted access (0.0.0.0/0) to specific ports. Unrestricted access increases opportunities for malicious activity (hacking, denial-of-service attacks, loss of data). The ports with highest risk are flagged red, and those with less risk are flagged yellow. Ports flagged green are typically used by applications that require unrestricted access, such as HTTP and SMTP.
- Security Groups - Unrestricted Access: Checks security groups for rules that allow unrestricted access to a resource. Unrestricted access increases opportunities for malicious activity (hacking, denial-of-service attacks, loss of data).
- IAM Use (Free!): Checks for your use of AWS Identity and Access Management (IAM).

- **Amazon S3 Bucket Permissions:** Checks buckets in Amazon Simple Storage Service (Amazon S3) that have open access permissions. This check examines explicit bucket permissions, but it does not examine associated bucket policies that might override the bucket permissions.
- **MFA on Root Account:** Checks the root account and warns if multi-factor authentication (MFA) is not enabled.
- **IAM Password Policy:** Checks the password policy for your account and warns when a password policy is not enabled or if password content requirements have not been enabled.
- **Amazon RDS Security Group Access Risk:** Checks security group configurations for Amazon Relational Database Service (Amazon RDS) and warns when a security group rule might grant overly permissive access to your database.

Fault Tolerance

- Amazon EBS Snapshots
- Load balancer optimization
- Auto Scaling Group Resources
- Amazon RDS Multi-AZ
- Amazon Route 53 name server delegations
- ELB connection draining



Fault Tolerance



9 ✓ 2 ▲

2 !

1 excluded items



20

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

AWS Trusted Advisor helps you increase the availability and redundancy of your AWS application by taking advantage of autoscaling, health checks, multi-AZ, and backup capabilities.

The following fault tolerance checks are available with Trusted Advisor:

- Amazon EBS Snapshots: Checks the age of the snapshots for your Amazon Elastic Block Store (Amazon EBS) volumes (available or in use).
- Load Balancer Optimization: Checks your load balancer configuration.
- Auto Scaling Group Resources: Checks the availability of resources associated with launch configurations and your Auto Scaling groups.
- Amazon RDS Multi-AZ: Checks for DB instances that are deployed in a single Availability Zone.
- Amazon Route 53 Name Server Delegations: Checks for Amazon Route 53 hosted zones for which your domain registrar or DNS is not using the correct Route 53 name servers.

- ELB Connection Draining: Checks for load balancers that do not have connection draining enabled.

Performance Improvement

- High-utilization Amazon EC2 instances
- Service limits
- Large number of rules in EC2 security group
- Over-utilized Amazon EBS magnetic volumes
- Amazon EC2 to EBS throughput optimization
- Amazon CloudFront alternate domain names



Performance



8 ✓ 0 ▲

0 !

0 excluded items



Training and

21

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

AWS Trusted Advisor helps you improve the performance of your service by checking your service limits, ensuring that you take advantage of provisioned throughput, and monitoring for over-utilized instances.

The following performance improvement checks are available with Trusted Advisor:

- High Utilization Amazon EC2 Instances: Checks the Amazon Elastic Compute Cloud (Amazon EC2) instances that were running at any time during the last 14 days and alerts you if the daily CPU utilization was more than 90% on four or more days.
- Service Limits: Checks for usage that is more than 80% of the service limit.
- Large Number of Rules in EC2 Security Group: Checks each Amazon EC2 security group for an excessive number of rules.
- Over-utilized Amazon EBS Magnetic Volumes: Checks for Amazon Elastic Block Store (EBS) Magnetic volumes that are potentially over-utilized and might benefit from a more efficient configuration.

- Amazon EC2 to EBS Throughput Optimization: Checks for Amazon EBS volumes whose performance might be affected by the maximum throughput capability of the Amazon EC2 instance they are attached to.
- CloudFront Alternate Domain Names: Checks Amazon CloudFront distributions for alternate domain names with incorrectly configured DNS settings.

Knowledge Check

True or False: Auto Scaling helps you ensure that you have the correct number of EC2 instances available to handle the load for your application.

True

Q: What feature would you use with an auto scaling policy to determine when your auto scaling group should scale out/in?

Amazon CloudWatch alarms

Q: You have an application composed of individual services and need to route a request to a service based on the content of the request. What type of load balancer should you use?

Application Load Balancer

Q: Which AWS service serves as a best practice and recommendation engine?

AWS Trusted Advisor

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



22


The answers to the knowledge check are shown in the slide.

Lab 3

Scale and Load Balance your Architecture

45 minutes

What You're Starting With – Lab 2 Review



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

24

Lab 3 Overview



- Auto Scaling
- **Create an AMI** from the running web server instance
- Add a **load balancer**
- Create a **launch configuration**
- Create an **auto scaling group**
- **Launch instances** within private subnet
- Add instances to load balancer

- **Test** auto scaling
- **Monitor** performance with CloudWatch alarms


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



Training and
Certification

25

Lab 3 – Scale and Load Balance your Architecture




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

26

Lab 3 – Scale and Load Balance your Architecture




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

27

Lab 3 – Scale and Load Balance your Architecture




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

28

Lab 3 – Scale and Load Balance your Architecture




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

29

Lab 3 – Scale and Load Balance your Architecture




© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
Training and
Certification

30

Module 6


Course Wrap-Up



The slide shows the learning path of concepts we have covered today.

Expand Your Cloud Skills with AWS


Online videos and labs



Start working with an AWS service in minutes with free online instructional videos and labs

[aws.amazon.com/training/
self-paced-labs](https://aws.amazon.com/training/self-paced-labs)


Instructor-led courses



Learn how to design, deploy, and operate highly available, cost-effective, and secure applications on AWS

aws.amazon.com/training

Certification



Validate your proven technical expertise with the AWS platform and gain recognition for your skills

aws.amazon.com/certification

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



3

AWS Training and Certification is an organization dedicated to expanding and deepening knowledge of AWS, as well as driving proliferation in the usage of AWS services. Our programs are designed for customers, partners and AWS employees.

We are continuously rolling out new and updated courses, training labs, and certifications to our customers and partners.

Self-Paced Labs

- Learn an individual [AWS Service topic](#)
- Follow a Learning Quest by [AWS Service Area or Use Case](#)
- Practice working with AWS as you [prepare for an exam](#)



For more information, see aws.amazon.com/training/self-paced-labs/.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and
Certification

4

Designed by AWS subject matter experts, self-paced labs provide an opportunity to use the AWS console with step-by-step instructions, giving you hands-on practice to help you gain confidence working with AWS.

AWS ILT Training Courses

	AWS Technical Essentials 1 day		
Introductory courses			
Intermediate courses	Architecting on AWS 3 days	Developing on AWS 3 days	Systems Operations on AWS 3 days
Advanced courses	Advanced Architecting on AWS 3 days	DevOps Engineering on AWS 3 days	Security Operations on AWS 3 days
Specialty courses	Migrating to AWS 2 days	Big Data on AWS 3 days	Data Warehousing on AWS 3 days

<https://aws.amazon.com/training/>

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification

The slide shows the current instructor-led training in the AWS Training & Certification catalog.

Classes are role-based and designed to be progressive as individuals gain deeper experience working with AWS services.

AWS Technical Essentials is our one-day foundational training on AWS services.

Architecting on AWS is a three-day course designed to teach Solutions Architects to design and use AWS services for common IT applications.

Developing on AWS is a three-day course designed to help individuals design and build secure, reliable and scalable AWS-based applications.

Systems Operations on AWS is a three-day course designed to help individuals operate highly available and scalable infrastructure on the AWS platform.

Advanced Architecting on AWS is a three-day course that covers building more complex solutions which incorporate data services, infrastructure configuration management, and security on AWS.

DevOps Engineering on AWS is a three-day advanced level course that demonstrates how to use the most common DevOps patterns to develop, deploy, and maintain applications on AWS.

Security Operations on AWS is a three-day advanced level course that demonstrates how to efficiently use AWS security services to stay secure and compliant.


Migrating to AWS is a two-day ILT that demonstrates cloud migration strategies, and dives into our recommended five phase plan for migrating existing workloads to the AWS Cloud.

Big Data on AWS is a three-day specialty course that introduces you to cloud-based big data solutions and Amazon Elastic MapReduce (EMR), the AWS big data platform.


Data Warehousing on AWS is a three-day specialty course designed to teach you how to design a cloud-based data warehousing solution using Amazon RedShift.

<https://aws.amazon.com/training/>


AWS Certification



AWS Certified Solutions Architect - Associate



AWS Certified Developer - Associate



AWS Certified SysOps Administrator- Associate

AWS Certified Solutions Architect - Professional

AWS Certified DevOps Engineer - Professional

For more information, see aws.amazon.com/certification.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



6

AWS certifications validate technical skills and knowledge necessary for designing, deploying, and operating solutions on the AWS platform. Earning certification helps you gain credibility for your proven experience working with AWS, and contributes to your organization's proficiency with AWS-based applications.

Benefits of AWS Certification

Individual

- Demonstrate expertise
- Stand out
- Industry visibility
- Customer visibility
- Peer recognition
- Credibility with customers

Employer

- Baseline bar on AWS skills
- Identify expert talent
- Leverage best practices
 - Reduce operational risk
 - Increase business advantage
 - Maximize AWS efficiencies
- Common vocabulary
- Accelerate time to cloud

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



7

Preparing for AWS Certification

For resources to help you prepare for the certification exam, see [aws.amazon.com/certification.](https://aws.amazon.com/certification/)

- [Exam Guides & Sample Questions](#)
- [AWS-Authored Study Guide](#)
- [Self-Paced Labs on **qwikLABS**](#)

[AWS Technical Training](#)

[AWS Whitepapers & FAQs](#)

[AWS Documentation & Reference Architectures](#)

[Practice Exams](#)

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



8

For more information, see:

- Prepare for Certification: <https://aws.amazon.com/certification/certification-prep/>
- Exam Guides: <https://aws.amazon.com/certification/certification-prep/#exam-guide-1>
- Order the AWS Authored Study Guide:
<https://www.amazon.com/dp/1119138558>
- Self-Paced Labs: <https://aws.amazon.com/certification/certification-prep/#exam-guide-1>
- AWS Technical Training: <https://aws.amazon.com/certification/certification-prep/#technical-training-1>
- AWS Whitepapers: <https://aws.amazon.com/whitepapers/>
- Practice Exams: <https://webassessor.com/aws/>

AWS Support

Support Options



The Technical Account Manager provides...

- ✓ A dedicated **voice within AWS** to serve as your **advocate**.
- ✓ **Proactive guidance** and **insight** into ways to optimize AWS through business and performance reviews.
- ✓ Orchestration and access to the full **breadth and depth of technical expertise** across the full range of AWS.
- ✓ Access to resources and **best practice recommendations**.



Infrastructure Event Management provides...

- ✓ A common understanding of event objectives and use cases through **pre-event planning and preparation**.
- ✓ Resource **recommendations** and deployment **guidance** based on anticipated capacity needs.
- ✓ **Dedicated attention** of the your AWS Support team during your event.
- ✓ The ability to immediately **scale down resources** to normal operating levels post-event.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

10

Your Technical Account Manger (TAM) is your primary point of contact who provides guidance and advocacy to proactively optimize your environment. The TAM sets the suite of AWS Support offerings and services in motion to prevent issues from arising.

For more information, see: <https://aws.amazon.com/premiumsupport/>

Support Options



The Concierge Service provides...

- ✓ A primary contact to help **manage AWS resources**.
- ✓ **Personalized handling** of billing inquiries, tax questions, service limits, and bulk reserve instance purchases.
- ✓ Direct access to an agent to help **optimize costs**, and identify **underutilized resources**.



AWS Trusted Advisor provides...

- ✓ Insight into how and where you can get the **most impact for your AWS spend**.
- ✓ Opportunities to **reduce your monthly spend** and retain or **increase productivity**.
- ✓ Guidance on getting the **optimal performance and availability** based on your requirements.
- ✓ Confidence that your environment is **secure**.

Cost Optimization



0 0 0

Performance



1 0 0

Security



0 2 0

Fault Tolerance



0 0 0

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Training and Certification

11

The Concierge Agent is a senior customer service professional who is familiar with you and your environment and is able to quickly and efficiently address billing and account inquiries, thus freeing up your time to run your business.

AWS Trusted Advisor is an online resource to help you reduce cost, increase performance and fault tolerance, and improve security by optimizing your AWS environment. Developed using best practices, Trusted Advisor provides real-time guidance for specific services.

Support Comparison

	Enterprise	Business	Developer	Basic
Customer Service 24x7x365	✓	✓	✓	✓
Support Forums	✓	✓	✓	✓
Documentation, White Papers, Best Practice Guides	✓	✓	✓	✓
AWS Trusted Advisor	Full Checks	Full Checks	Basic Checks	Basic Checks
Access to Technical Support	Phone, chat, email, live screen sharing, TAM (24/7)	Phone, chat, email, live screen sharing	Email (local business hours)	Support for Health Checks
Primary Case Handling	Sr. Cloud Support Engineer	Cloud Support Engineer	Cloud Support Associate	Technical Customer Service Associate
Users who can create Technical Support cases	Unlimited (IAM supported)	Unlimited (IAM supported)	1 (account credentials only)	
Case Severity/Response Times	Critical: < 15 minutes Urgent: < 1 hour High: < 4 hours Normal: < 12 hours Low: < 24 hours	Urgent: < 1 hour High: < 4 hours Normal: < 12 hours Low: < 24 hours	Normal: < 12 hours Low: < 24 hours	
Architecture Support	Application Architecture	Use case guidance	Building blocks	
Best Practice Guidance	✓	✓	✓	
Client-side Diagnostic Tools	✓	✓	✓	
AWS Support API	✓	✓	✓	
Third-Party Software Support	✓	✓	✓	
Infrastructure Event Management	✓	Available at additional cost		
AWS Concierge	✓			
Direct access to Technical Account Manager (TAM)	✓			
Prioritized Case Routing	✓			
Management Business Reviews	✓			

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification 12

For more information, see: <https://aws.amazon.com/premiumsupport/compare-plans/>

Module 7

Course Appendix

Module 1 Appendix

AWS Introduction and History



2

Cloud Computing Concepts



3

Understand essential characteristics of cloud computing.

What is cloud computing?

Cloud computing is on-demand delivery of IT resources and applications via the Internet with pay-as-you-go pricing.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



4


Cloud computing is a common term for a variety of computing concepts that involve large numbers of computers that are connected through a real-time communication network, like the Internet.

Cloud computing is the use of computing resources (hardware and software) that are delivered as a service over a network (typically the Internet). The name comes from the use of a cloud-shaped symbol as an abstraction in system diagrams for the complex infrastructure it contains. Cloud computing entrusts remote services with a user's data, software, and computation. Cloud computing allows you to access as many resources as you need, almost instantly, and only pay for what you use.

Instead of buying, owning, and maintaining your own data centers and servers, organizations can acquire technology such as compute power, storage, databases, and other services on an as-needed basis.

Word cloud created at: <http://www.tagxedo.com/>

Essential Characteristics of Cloud Computing



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.





5

Cloud computing is characterized by five features: on-demand self services, broad network access, resource pooling, rapid elasticity, and measured service.

On-Demand Self Services & Broad Network Access

- User provisions computing resources as needed.
- User interacts with cloud service provider through an online control panel.
- Clear solutions are available through a variety of network-connected devices and over varying platforms.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

6

On-demand self services are an essential characteristic of cloud computing. The user provisions computing resources as needed and interacts with the cloud service provider through an online control panel, such as the AWS Management Console.

Broad network access is another characteristic of cloud computing. Clear solutions are available through a variety of network-connected devices and over varying platforms.

Resource Pooling

Securely separate resources to service multiple customers.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon
web services

Training and
Certification

7

A third characteristic of cloud computing is resource pooling. Cloud computing solutions securely separate resources to service multiple customers.

Rapid Elasticity

Resources are quickly scalable and flexible based on business needs.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon
web services

Training and
Certification

8

Customers used to over-provision to ensure that they had enough capacity to handle their business operations at the peak level of activity. Cloud computing allows customers to provision the number of resources that they actually need, knowing they can instantly scale up or down along with the needs of their business, which also reduces cost and improves the customer's ability to meet their users' demands.

Measured Service

Pay for services as you go.

Electrical services
analogy



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services


Training and
Certification

9

Pay for service you use as you go.

The slide shows an analogy that helps explain cloud computing. Electricity services are a utility that you pay for on-demand: you pay for what you use. You plug electrical appliances into a vast electrical grid that is managed by the power company to get a low-cost, reliable power supply. This power is available to you from the power company with much greater efficiency than you could generate on your own.

What Does My AWS Cloud Look Like?



amazon
Training and
Certification

10

The slide shows the architecture of a highly scalable and reliable web application on AWS. This is an example of one possible configuration among millions.

For more information, see: <http://docs.aws.amazon.com/gettingstarted/latest/wah-linux/web-app-hosting-intro.html>


Module 2 Appendix

AWS Foundational Services



11

Data Center Design Models



The one-tier model applies to clients (no local processing or storage) connected to a mainframe. This design is usually used for monolithic applications such as kiosks. This model has limited scalability and lacks flexibility.

The two-tier model applies to clients connected to database servers. The clients have direct interaction with the database server and do some local application processes. This model has limited scalability and is not recommended for critical applications.

The three-tier model applies to clients connected to application servers. The application servers are then connected to the database servers. This model has more scalability than the first two models.

The n-tier model consists of clients connected to n number of application servers, connected to n number of database servers. This model has the most scalability of traditional data center design models and a robust partitioning of application functionality.

Web Services Model



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services

Training and
Certification

14

The Web Services Model includes web server, application server, and database server tiers. It includes tasks performed by multiple hosts with specific roles.

Amazon EC2

AMI Types - Storage for the Root Device



Characteristic	Amazon EBS-Backed	Amazon Instance Store-Backed
Boot time	Usually < 1 minute	Usually < 5 minutes
Size limit	16 TiB	10 GiB
Data persistence	The root volume is deleted when the instance terminates. Data on any other Amazon EBS volumes persists after the instance is terminated.	Data on any instance store volumes persists only during the life of the instance.
Charges	Instance usage, Amazon EBS volume usage, and storing your AMI as an Amazon EBS snapshot.	Instance usage and storing your AMI in Amazon S3.
Stopped state	Can be stopped.	Cannot be stopped.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Training and Certification

16

AMIs are either Amazon Elastic Block Storage (EBS)-backed or backed by instance store. When an AMI is EBS-backed, the root device for an instance is an EBS volume created from an EBS snapshot. When an AMI is instance-store backed, the root device for the instance was created from a template stored in Amazon S3. Key differences between the two categories of AMIs are shown in the slide.


Storage Concepts and Solutions




17

Understand common storage concepts and solutions related to servers and application environments.

Block and File Level Storage



Block Level Storage



File Level Storage


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon web services

18

Block-Level Storage: This is the reading and writing of data to and from a disk using a sequence of bytes in a predetermined length. Files are spread over many blocks. Block-level incremental backups are faster because they do not have to back up the entire file each time. Instead they only backup new blocks or blocks that have changed since the last backup. Example of block storage solutions are Direct-Attached Storage (DAS) and network-based Storage Area Network (SAN). Common DAS protocols are ATA, SATA, SCSI, SAS, and USB. Common SAN protocols are iSCSI and FCoE. With block-level storage, a logical unit number (LUN) can be treated as a physical drive. Files are divided into many blocks and if any identical blocks are found, redundant copies are eliminated.

File-Level Storage: In file-level storage, files and folders can be accessed and managed by the storage system, but the smaller storage blocks that make up the files and folders cannot be directly controlled. Storage drives need to be configured with a storage protocol like Network File System (NFS), Server Message Block (SMB), or Common Internet File System (CIFS). With file-level storage, a shared folder can be mounted as a network drive. If one file has to be accessed on different systems, file level storage is ideal choice.



The slide shows commonly used storage technologies.

Direct-Attached Storage (DAS): DAS is storage directly attached to the computer or client using it.

Network-Attached Storage (NAS) Systems: NAS appliances are servers configured with software designed specifically for storing and providing access to files over a LAN.

Storage-Attached Network (SAN) Architecture: A SAN architecture is composed of servers in a network connected to centralized disk storage.

Amazon S3

Amazon S3 Buckets



- Organize the Amazon S3 namespace at the highest level.
- Identify the account responsible for storage and data transfer charges.
- Play a role in access control.
- Serve as the unit of aggregation for usage reporting.
- Have globally unique bucket names, regardless of the AWS region in which they were created.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 amazon
web services

Training and
Certification

21

A bucket is a logical container for objects stored in Amazon S3. Every object is contained in a bucket. Buckets serve several purposes: They organize the Amazon S3 namespace at the highest level, they identify the account responsible for storage and data transfer charges, they play a role in access control, and they serve as the unit of aggregation for usage reporting. Amazon S3 bucket names are globally unique, regardless of the AWS Region in which you create the bucket. You specify the name at the time you create the bucket.

Amazon S3 Region Considerations



- Amazon S3 creates a bucket in the region you select.
- You can choose a region to:
 - Optimize latency
 - Minimize costs
 - Address regulatory requirements
- Objects stored in a region never leave the region unless you explicitly transfer them to another region.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services

Training and
Certification

22

For more information, see:

http://docs.aws.amazon.com/general/latest/gr/rande.html#s3_region

Amazon S3 Objects



- Objects are the fundamental entities stored in Amazon S3.
- When using the console, you can think of them as files.
- **Objects consist of data and metadata.** The data portion is opaque to Amazon S3. The metadata is a set of name-value pairs that describe the object.
 - Default metadata such as the date last modified
 - Standard HTTP metadata such as Content-Type
 - Custom metadata at the time the object is stored
 - A key that uniquely identifies an object within its bucket

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services

Training and
Certification


23

Objects are the fundamental entities stored in Amazon S3. When using the console, you can think of them as being files. Objects consist of data and metadata. The data portion is opaque to Amazon S3. The metadata is a set of name-value pairs that describe the object. These include some default metadata, such as the date last modified, and standard HTTP metadata such as Content-Type. You can also specify custom metadata at the time the object is stored. An object is uniquely identified within a bucket by a key.

For more information, see: <http://docs.aws.amazon.com/AmazonS3/latest/dev/s3-dg.pdf>

Amazon S3 + Amazon Glacier

S3 Lifecycle policies allow you to delete or move objects based on age and set rules per S3 bucket.



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

24

Example:

1. Move object to Amazon Glacier after 30 days
2. Delete object after 365 days

For more information, see:

<http://docs.aws.amazon.com/AmazonS3/latest/dev/object-lifecycle-mgmt.html>

Amazon EBS

EBS Performance



EBS Magnetic

- 40-200 IOPS

EBS General Purpose SSD

- SSD backed
- 3 IOPS / GB
- Burstable to 3,000 IOPS and up to 10,000 IOPS

EBS Provisioned IOPS SSD

- SSD backed
- Up to 20,000 IOPS consistently
- Up to 320 MB/s throughput

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon EBS volume types are shown in the slide.

For more information, see:

<http://docs.aws.amazon.com/AWSEC2/latest/UserGuide/EBSVolumeTypes.html>

Amazon CloudFront

Amazon CloudFront



Amazon
CloudFront


- Easy and cost effective way to **distribute content** to end users
- **Low latency, high data transfer speeds**
- Deliver your entire website, including static, dynamic, and streaming content using a global network of edge locations

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.


Training and
Certification

28

Amazon CloudFront integrates with other AWS services to give developers and businesses an easy way to distribute content to end users with low latency, high data transfer speeds, and no minimum commitments. Amazon CloudFront is a web service that speeds up distribution of your static and dynamic web content, for example, .html, .css, .php, and image files, to end users. Amazon CloudFront delivers your content to edge locations.



1. Configure your origin servers.
2. Upload files to origin servers.
3. Create an Amazon CloudFront distribution.
4. Amazon CloudFront sends your distribution's configuration to all of its edge locations.
5. As you develop your website or application, use the domain name that Amazon CloudFront provides for your URLs.
6. Optionally, you can configure your origin server to add expiry headers to the files; the headers indicate how long you want the files to stay in the cache in Amazon CloudFront edge locations.



1. A user accesses your website or application and requests one or more objects, such as an image file.
2. DNS routes the request to the Amazon CloudFront edge location that can best serve the user's request - typically the nearest CloudFront edge location in terms of latency.
3. In the edge location, CloudFront checks its cache for the requested files. If the files are in the cache, CloudFront returns them to the user. If the files are not in the cache, it does the following:
 - a. CloudFront compares the request with the specifications in your distribution and forwards the request for the files to the applicable origin server for the corresponding file type—for example, to your Amazon S3 bucket for image files.
 - b. The origin servers send the files back to the CloudFront edge location.
 - c. As soon as the first byte arrives from the origin, CloudFront begins to forward the files to the user. CloudFront also adds the files to the cache in the edge location for the next time someone requests those files.
4. After an object has been in an edge cache for 24 hours or for the duration specified in your file headers, CloudFront does the following:
 - a. CloudFront forwards the next request for the object to your origin to determine whether the edge location has the latest version.

- b. If the version in the edge location is the latest, CloudFront delivers it to your user.
- c. If the version in the edge location is not the latest, your origin sends the latest version to CloudFront, and CloudFront delivers the object to your user and stores the latest version in the cache at that edge location.

Networking Concepts



31


Understand typical networking devices, protocols, and services.

What is a Network?

A network is two or more computers linked to share resources, exchange files, or allow electronic communications.

Network Types:

- Local Area Network (LAN)
- Wide Area Network (WAN)
- Virtual Private Network (VPN)



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon
web services

Training and
Certification

32

A local area network (LAN) covers a local area such as an office or home.

A wide area network (WAN) covers an area wider than a LAN. A WAN can range from a corporate network to the Internet.

A virtual private network (VPN) allows users to securely connect to a network through the Internet and remotely access network resources.

Physical vs. Logical Topology


- A physical topology defines how the systems are physically connected.
- A logical topology defines how the systems communicate across the physical topologies.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



33

It is important to understand the difference between a physical and logical topology. You should know how a network is laid out and how the devices communicate on that network to make security decisions to protect your environment.



The slide shows a basic network diagram consisting of various physical network devices.

Servers are very fast computers with a large amount of RAM and storage space and one or more fast network interface cards. Servers are the central repository of data and applications shared by users in a network.

Workstations are computers and devices with a network interface card or wireless adapter to allow quick connections to networks.

Switches are devices that provide a central connection point for cables from workstations, servers, and peripherals.

Routers forward data packets between computer networks.

Firewalls are hardware or software that help secure your network by creating rules to block/allow access.

Amazon VPC

Networking in Your VPC



You can use the following components to configure networking in your VPC:

- IP addresses
- Elastic network interfaces
- Route tables
- Internet gateways
- Network Address Translation (NAT)
- Dynamic Host Configuration Protocol (DHCP) options sets
- Domain Name System (DNS)
- VPC peering
- VPC endpoints
- VPC flow logs

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services

Training and
Certification

36

Flow logs capture information about the IP traffic going to and from network interfaces in your VPC.

The diagram shows the layers of security provided by security groups and network ACLs.

For more information, see:

- http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/VPC_Networking.html
- <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/flow-logs.html>

Module 3 Appendix

Security, Identity, and Access Management



Training and
Certification

37

Data Center Security



38

Understand common methods used to secure data centers.

Physical & Environmental Security



- Lock your data center.
- Only provide access to those who need it.
- Keep track of access.
- Mount servers on racks with locks.
- Have redundant utilities.
- Build your data center with security in mind.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
web services

Training and
Certification

39

Physical and environmental security is important in an on-premises data center. Remember to lock your data center, only provide access to people who need it, and keep track of the access. Servers should be mounted on racks with locks. Ideally, data centers should have redundant utilities (electricity, water, voice, data, HVAC).

Network Security

- Identification & Authentication
- Firewalls
- Patching
- Virus Protection
- Encryption

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



40

Establishing and maintaining a secure computing environment is more difficult when networks become more interconnected. The slide shows common methods used to security your network and ensure confidentiality, integrity, and availability of systems and data.

Identification and Authentication: The careful user of user accounts is an important aspect of security in your network. Least privilege access should be provided and user accounts should be properly maintained to avoid unauthorized access to resources.

Firewalls: A firewall is a security router that sits between the Internet and your network. The firewall's purpose is to act as a security guard between the Internet and the network it is guarding.

Patching: Maintaining operating system updates (patching) is another important aspect of network security. Software patches are minor updates that fix bugs and address security flaws that can be abused by hackers. Keeping your operating systems up to date with the latest patches can help facilitate a safer environment for your network.

Virus Protection: Antivirus programs can help detect and remove known threats to clients on your network.

Encryption: Encryption processes use numeric keys and algorithms to scramble data when it is sent over the network or saved on disk. This helps prevent unauthorized users from reading and accessing confidential or sensitive data. Effective encryption is important in the effective use of a virtual private network (VPN).

AWS IAM

Advanced Concepts

AWS Resource-Based Policies

- Are an alternative to IAM and supported by some services.
- Grant cross-account access to your resources.
- Use a principal to uniquely identify accounts in the policy.
- Supported AWS services include :
 - Amazon S3 Bucket Policy
 - Amazon SNS Topic Policy
 - Amazon SQS Queue Policy
 - Amazon Glacier Vault Policy
 - AWS OpsWorks Stack Policy
 - AWS Lambda Function Policy

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



42

For some AWS services, you can grant cross-account access to your resources. To do this, attach a policy directly to the resource that you want to share, instead of using a role as a proxy. The resource that you want to share must support resource-based policies. Unlike a user-based policy, a resource-based policy specifies who (in the form of a list of AWS account ID numbers) can access that resource. Cross-account access with a resource-based policy has an advantage over a role. With a resource that is accessed through a resource-based policy, the user still works in the trusted account and does not have to give up his or her user permissions in place of the role permissions. In other words, the user continues to have access to resources in the trusted account at the same time as he or she has access to the resource in the trusting account. This is useful for tasks such as copying information to or from the shared resource in the other account.

Principal: This element defines an account in a policy. In a resource-based policy, the principal may refer to the same account or another account.

For more information, see:

http://docs.aws.amazon.com/IAM/latest/UserGuide/id_roles_compare-resource-policies.html

Access to AWS Resources



Temporary Security Credentials

- Security Token Service
- AssumeRole
- AssumeRoleWithSAML
- AssumeRoleWithWebIdentity

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon
Training and
Certification

43

For more information, see:

- AssumeRole -
http://docs.aws.amazon.com/STS/latest/APIReference/API_AssumeRole.html
- GetFederationToken -
http://docs.aws.amazon.com/STS/latest/APIReference/API_GetFederationToken.html

AWS Services support for IAM Roles



- AWS CLI on Amazon EC2
- AWS CloudTrail logs to Amazon S3
- Amazon Elastic Transcoder access to Amazon S3
- AWS Elastic Beanstalk access to AWS services
- AWS Lambda code access to AWS services
- Many more ...

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

44

IAM roles may be used in many scenarios to allow AWS services and resources to gain access to resources in the same or another account. Some uses cases are mentioned in the slide.

Module 4 Appendix Databases

Security Groups



Allow access to IP address ranges or Amazon EC2 instances you specify.

Use VPC security groups to control access to a DB instance inside a VPC.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon
Training and
Certification

46

A VPC security group controls access to a DB instance inside a VPC.

DB Parameter & Option Groups



DB parameter groups:

- Contain engine configuration values that can be applied to one or more DB instances of the same instance type.
- Are applied by Amazon RDS by default when you create DB instance, which contains defaults for the specific database engine and instance class of the DB instance.

DB option groups:

- Tools that simplify database management.
- Currently available for Oracle, Microsoft SQL Server, and MySQL 5.6 DB instances.

Configuration Details

Engine:	sqlserver-web (11.00.2100.60.v1)
DB Name:	tempdb
Username:	sa
Option Group(s):	default.sqlserver-web-11-00 (in-sync)
Parameter Group:	sqlsvr-web11-parms (pending-reboot)

You manage the configuration of a DB engine by using a DB parameter group. A DB parameter group contains engine configuration values that can be applied to one or more DB instances of the same instance type. Amazon RDS applies a default DB parameter group if you don't specify a DB parameter group when you create a DB instance. The default group contains defaults for the specific database engine and instance class of the DB instance.

Some DB engines offer tools that simplify managing your databases and making the best use of your data. Amazon RDS makes such tools available through option groups.

For more information, see:

- Oracle Database Engine option groups -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.Oracle.Options.html>
- SQL Server option groups -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.SQLServer.Options.html>
- MySQL option groups -
<http://docs.aws.amazon.com/AmazonRDS/latest/UserGuide/Appendix.MySQL.Options.html>

Supported Operations



Table Operations:

- Create, update, and delete tables.
- After creation, you can increase or decrease provisioned throughput.
- Retrieve the table's status, the primary key, and when the table was created.
- List all tables in your account for a region.

Item Operations:

- Add, update, and delete items from a table.
- Add, update, and delete existing attributes from an item.
- Perform conditional updates.
- Retrieve a single item or multiple items.


© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 **amazon**
webservices

48

Amazon DynamoDB provides operations to create, update, and delete tables. After the table is created, you can use the UpdateTable operation to increase or decrease a table's provisioned throughput. Amazon DynamoDB also supports an operation to retrieve table information (the DescribeTable operation), including the current status of the table, the primary key, and when the table was created. The ListTables operation enables you to get a list of tables in your account in the region of the endpoint you are using to communicate with Amazon DynamoDB.


Item operations enable you to add, update, and delete items from a table. You can update existing attribute values, add new attributes, and delete existing attributes from an item. You can also perform conditional updates. For example, if you are updating a price value, you can set a condition so the update happens only if the current price is \$15. You can retrieve a single item or multiple items from a table.



If you want to read the data using non-key attributes, you can use a secondary index to do this. A local secondary index is an index that has the same partition key as the table, but a different sort key.

For more information, see:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/LSI.html>



A global secondary index is an index with a partition key and sort key that can be different from those on the table. They can be thought of as “pivot charts” for your table. The global secondary index must be created at the time the table is created.


For more information, see:

<http://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GSI.html>

Module 5 Appendix

AWS Elasticity and Management Tools


AutoScaling Advanced Concepts



For more information, see:

http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/scaling_plan.html#scaling_typesof


Elastic Load Balancing Advanced Concepts



Internet-facing load balancer: An Internet-facing load balancer takes requests from clients over the Internet and distributes them across the EC2 instances that are registered with the load balancer.

Internal load balancer: An internal load balancer routes traffic to your EC2 instances in private subnets. The clients must have access to the private subnets.

HTTPS load balancer: You can create a load balancer that uses the SSL/TLS protocol for encrypted connections (also known as SSL offload). This feature enables traffic encryption between your load balancer and the clients that initiate HTTPS sessions, and for connections between your load balancer and your backend instances.



Before a client sends a request to your load balancer, the client resolves the load balancer's domain name using a Domain Name System (DNS) server. The DNS entry is controlled by Amazon, because your instances are in the `amazonaws.com` domain. The Amazon DNS servers return one or more IP addresses to the client. These are the IP addresses of the load balancer nodes for your load balancer. As traffic to your application changes over time, Elastic Load Balancing scales your load balancer and updates the DNS entry. Note that the DNS entry also specifies the time-to-live (TTL) as 60 seconds, which ensures that the IP addresses can be remapped quickly in response to changing traffic.

The client uses DNS round robin to determine which IP address to use to send the request to the load balancer. The load balancer node that receives the request uses a routing algorithm to select a healthy instance. It uses the round robin routing algorithm for TCP listeners, and the least outstanding requests routing algorithm (favors the instances with the fewest outstanding requests) for HTTP and HTTPS listeners.

The cross-zone load balancing setting also determines how the load balancer selects an instance. If cross-zone load balancing is disabled, the load balancer node selects the instance from the same Availability Zone that it is in. If cross-zone load balancing is enabled, the load balancer node selects the instance regardless of Availability Zone. The load balancer node routes the client request to the selected instance.

Listeners



- A listener is a process that checks for connection requests.
- Front-end connections are:
 - Client to load balancer connections.
 - Configured with a protocol and a port.
- Back-end connections are:
 - Load balancer to back-end instance connections.
 - Configured with a protocol and a port .
- ELB supported protocols:
 - HTTP
 - HTTPS
 - TCP
 - SSL

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

Amazon Web Services | Training and Certification

57

Before you start using Elastic Load Balancing, you must configure one or more listeners for your load balancer. A listener is a process that checks for connection requests. It is configured with a protocol and a port for front-end (client to load balancer) connections, and a protocol and a port for back-end (load balancer to back-end instance) connections.

Elastic Load Balancing supports the following protocols:

- HTTP
- HTTPS (secure HTTP)
- TCP
- SSL (secure TCP)

The HTTPS protocol uses the SSL protocol to establish secure connections over the HTTP layer. You can also use the SSL protocol to establish secure connections over the TCP layer.

If the front-end connection uses TCP or SSL, your back-end connections can use either TCP or SSL. If the front-end connection uses HTTP or HTTPS, your back-end connections can use either HTTP or HTTPS.

Back-end Instances for Your Load Balancer



- Health checks
- Security groups
- Subnets
- Register
- De-register instances

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon web services

Training and
Certification

58

After you've created your load balancer, you must register your EC2 instances with the load balancer. You can select EC2 instances from a single Availability Zone or multiple Availability Zones within the same region as the load balancer. Elastic Load Balancing routinely performs health checks on registered EC2 instances and automatically distributes incoming requests to the DNS name of your load balancer across the registered, healthy EC2 instances.

Health checks are periodic pings, attempted connections, or requests sent to EC2 instances by your load balancer to check the availability of your EC2 instances. The load balancer performs health checks on all registered instances, whether the instance is in a healthy state or an unhealthy state. The load balancer routes requests only to the healthy instances. When the load balancer determines that an instance is unhealthy, it stops routing requests to that instance. The load balancer resumes routing requests to the instance when it has been restored to a healthy state.

A security group acts as a firewall that controls the traffic allowed to and from one or more instances. When you launch an EC2 instance, you can associate one or more security groups with the instance. For each security group, you add one or more rules to allow traffic. You can modify the rules for a security group at any time; the new rules are automatically applied to all instances associated with the security group. You must ensure that the security groups for your instances allow the load balancer to communicate with your back-end instances on both the listener port and the health check port. In a VPC, your security groups and network access control lists (ACL) must allow traffic in both directions on these ports.

When you attach a subnet to your load balancer, Elastic Load Balancing creates a load balancer node in the Availability Zone. Load balancer nodes accept traffic from clients and forward requests to the healthy registered instances in one or more Availability Zones. For load balancers in a VPC, we recommend that you attach one subnet per Availability Zone for at least two Availability Zones. This improves the availability of your load balancer. Note that you can modify the subnets attached to your load balancer at any time.

Registering an EC2 instance adds it to your load balancer. The load balancer continuously monitors the health of its registered instances and routes requests to the healthy registered instances. If demand on your instances increases, you can register additional instances with the load balancer to handle the demand.

De-registering an EC2 instance removes it from your load balancer. The load balancer stops routing requests to an instance as soon as it is de-registered. If demand decreases, or you need to service your instances, you can de-register instances from the load balancer. A de-registered instance remains running but no longer receives traffic from the load balancer, and you can register it with the load balancer again when you are ready.

For more information, see:

<http://docs.aws.amazon.com/ElasticLoadBalancing/latest/DeveloperGuide/elb-backend-instances.html>

CloudWatch Advanced Concepts

CloudWatch Alarms

After 3 periods over threshold, an action is invoked

Only one period over threshold, no action is invoked

Threshold

Value

Units


Time Periods

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

You can create a CloudWatch alarm that sends an Amazon Simple Notification Service (SNS) message when the alarm changes state. An alarm watches a single metric over a time period you specify and performs one or more actions based on the value of the metric relative to a given threshold over a number of time periods. The action is a notification sent to an Amazon SNS topic or Auto Scaling policy. Alarms invoke actions for sustained state changes only. CloudWatch alarms will not invoke actions simply because they are in a particular state: the state must have changed and been maintained for a specified number of periods.

In the slide, the alarm threshold is set to 3 and the minimum breach is three periods. The alarm invokes its action only when the threshold is breached for three consecutive periods. In the figure, this happens with the third through fifth time periods, and the alarm is triggered. At period six, the value dips below the threshold, and the state is set to OK. Later, during the ninth time period, the threshold is breached again, but not for the necessary three consecutive periods. Consequently, the alarm's state remains OK.



The slides shows AWS services that Amazon CloudWatch collects metrics from.

For more information, see:

http://docs.aws.amazon.com/AmazonCloudWatch/latest/DeveloperGuide/supported_services.html

Module 6 Appendix

Course Wrap-Up

AWS Support

Case Severity & Response Times

	Critical	Urgent	High	Normal	Low
Enterprise Plan (24 x 7)	15 minutes or less	1 hour or less	4 hours or less	12 hours or less	24 hours or less
Business Plan (24 x 7)		1 hour or less	4 hours or less	12 hours or less	24 hours or less
Developer Plan (Business hours)				12 hours or less	24 hours or less

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



64

Critical: Your business is at risk. Critical functions of your application are unavailable.

Urgent: Your business is significantly impacted. Important functions of your application are unavailable.

High: Important functions of your application are impaired or degraded.

Normal: Non-critical functions of your application are behaving abnormally, or you have a time-sensitive development question.

Low: You have a general development question, or you want to request a feature.

Pricing

Basic	Developer	Business	Enterprise
Included	\$29/month -or- 3% of monthly AWS spend	Greater of \$100 -or- 10% of monthly AWS usage for the first \$0-\$10K 7% of monthly AWS usage from \$10K-\$80K 5% of monthly AWS usage from \$80K-\$250K 3% of monthly AWS usage over \$250K	Greater of \$15,000 -or- 10% of monthly AWS usage for the first \$0-\$150K 7% of monthly AWS usage from \$150K-\$500K 5% of monthly AWS usage from \$500k-\$1M 3% of monthly AWS usage over \$1M

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



65

All AWS Support tiers include an unlimited number of support cases, with no long-term contracts. Also, with the Business and Enterprise-level tiers, as your AWS charges grow, you earn volume discounts on your AWS Support costs.

For more information, see: <http://calculator.s3.amazonaws.com/index.html>

Pricing Examples

Business Pricing Example

For \$85K in AWS monthly usage:

$$\$10,000 \times 10\% = \$1,000$$

(10% of the first \$0 - \$10K of usage)

$$+ \$70,000 \times 7\% = \$4,900$$

(7% of usage from \$10K - \$80K)

$$+ \$5,000 \times 5\% = \$250$$

(5% of usage from \$80K - \$250K)

$$+ \$0 \times 3\% = \$0$$

(3% of usage over \$250K)

Total: \$6,500

Enterprise Pricing Example

For \$1.2M in AWS monthly usage:

$$\$150,000 \times 10\% = \$15,000$$

(10% of the first \$0 - \$150K of usage)

$$+ \$350,000 \times 7\% = \$24,500$$

(7% of usage from \$150K - \$500K)

$$+ \$500,000 \times 5\% = \$25,000$$

(5% of usage from \$500K - \$1M)

$$+ \$200,000 \times 3\% = \$6,000$$

(3% of usage over \$1M)

Total: \$70,500

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



66

Examples of business and enterprise pricing are shown on the slide.

Module 8

ILT- Specific Information

Course Feedback

Course Feedback

- In a web browser, navigate to
 - Customers: <https://awstraining.csod.com>
 - Partners: <https://www.apn-portal.com>
- Log in using the same credentials as your class registration.
- Point to My Training and click View Your Transcript.
- If you don't see AWS Technical Essentials on the Active list, click Active and then click Completed.
- Click the drop-down box and select Evaluate.

The screenshot shows the AWS Training and Certification portal interface. At the top, there is a navigation bar with links for Home, Calendar, My Training (which is highlighted with a red box), and ILT. Below the navigation bar, a greeting says "Hello [REDACTED] welcome 1" and a link to "View Your Transcript". A dropdown menu for filtering training items is open, showing options: Active (with a checked checkbox), Completed (which is also highlighted with a red box), Archived, and another Completed option. Below this, a specific course entry for "AWS Technical Essentials" is displayed, showing it starts on [REDACTED], has no due date, and is in a registered status. To the right of the course entry is a dropdown menu with options: View Training Details (which is highlighted with a red box), Manage, Evaluate (which is also highlighted with a red box), and another View Training Details option. At the bottom of the page, there is a copyright notice: "© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved." and the Amazon logo.

You will also receive an email with a link to the course evaluation upon completion of the course.

Gilmore & qwikLABS Access Student Guides, Lab Guides, Lab Environment




4

Student and Lab Guides - Gilmore

1 Sign in <http://online.vitalsource.com>

2 Enter **License Code** in email\handout

- Order Number:1017806
- Order Date:06/29/2015
- Product Name: [REDACTED]
- License Code: SVHGTCYQ6ZCBTS7MND8Q
- License Quantity: 1
- License Expiration Date: n/a
- Comments: Student Guide



© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

 Amazon web services | Training and Certification

5

QwikLABS Access

Sign in (Safari, Chrome, or Firefox preferred)
<https://aws.qwiklab.com>

The screenshot shows the QwikLABS access interface. On the left, under 'Class Details', there are two items: '- Lab 1 - Working with' and '- Lab 2 - Creating'. On the right, a specific lab is selected: '- Lab 1 - Working'. This lab card includes the Amazon logo, a brief description ('In this lab, you'll explore some of the key features of Amazon'), and details like Duration (60 min.), Access Time (120 min.), Setup Time (1 min.), and Level. A red arrow points from the 'Select' button in the lab card to the 'Start Lab' button at the bottom. The 'Start Lab' button is highlighted with a green box. At the bottom of the page, there is a copyright notice and the Amazon Web Services logo.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

Signing in to the AWS Management Console from QwikLABS

The screenshot shows two side-by-side interfaces. On the left, a 'CONNECTION' interface has fields for 'AWS Account' (112233445566), 'User Name' (awsstudent), and 'Password' (cgNb8GmXm4), with a blue 'Open Console' button. On the right, an AWS sign-in interface has fields for 'Account' (112233445566), 'User Name' (awsstudent), and 'Password' (*****), with a blue 'Sign In' button. Orange arrows point from the 'User Name' and 'Password' fields in the connection interface to their counterparts in the sign-in interface.

Gilmore VitalSource Bookshelf contains the lab manual.

NOTE: Do not change the default selected region.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

amazon web services | Training and Certification

7

Lab Sign In – Command Reference File

Lab 1 - Command Reference File CONNECTION ADDL. INFO

Lab 1 - Command Reference File

 [lab1-command-reference-file.txt](#)

Download the file to correctly copy the latest scripts used in the lab.

NOTE: This file is not a lab guide.

© 2017, Amazon Web Services, Inc. or its Affiliates. All rights reserved.



8